Referee-facing render of the final merged manuscript. Math typeset with MathJax. Provenance/build comments in the source are non-rendering. PREPARED DRAFT — not submitted.

13D Four-Force Interface from the Constraint-Selected GUT Geometry

Paper II of the Fable pair — restructured build · merged, reader-optimized statement, 2026-06-13

The four forces are the symmetries of one small frozen shape.

Take a single six-dimensional shape, freeze it, and let the four forces be its symmetries — with zero new dials between the five strengths, two of them (e and G_F) not predictions but identities that break against measurement if the geometry is wrong (§1.2.1). It is sharp enough to be wrong on a napkin, and it tells you exactly where.

THE BIG IDEA (in ~150 words)

Textbooks hand you four forces with five separate dial-like strengths and five disjoint origin stories: gravity with its own G_N, the strong force with g_s, the weak force with g, electromagnetism with e, and the Fermi scale G_F — five disjoint origins, five effective theories that do not talk to each other.

This paper makes one bet instead. Take a single small internal shape — the six-dimensional flag manifold $K_6 = SU(3)/T^2$, frozen by the companion GUT manuscript — and *let the forces be its symmetries*. The strong force is the shape's surviving rotation symmetry; the weak and electromagnetic forces descend from the two companion factors riding alongside it. Four explicit maps route each force out of the one shape. The payoff is thrift: the five force strengths add **zero new dials** — each is either read off the one frozen shape or computed from the others. Two of them, e and G_F, aren't predictions at all but identities — and an identity that disagreed with experiment would prove the geometry wrong (§1.2.1). Then an independent reality-check that the geometry is *buildable*: the same coset, built and simulated on open tools, is the admissibility chamber of a working error-corrector (§16, simulated grade — it changes no physics status). *Same shape, four forces, one machine — and it tells you exactly where it would break.*

Provenance. This is the four-force interface companion to Paper I (*Fable_GUT*), routing gravity and the strong, weak, and electromagnetic sectors through one frozen 13D geometry as four mutually compatible projections, with its own machine-certificate layer (Appendix Q). It re-closes none of Paper I's gates; all gate certificates remain external (§14). The reader's entry point is Section 0 (Reviewer First Read).

WHY THIS MATTERS (read this first)

One geometry, asked to source all four forces at once — with zero dials of its own — and it is sharp enough to be wrong on a napkin.

The stake is not "a parent group exists." It is concrete: **here are the four explicit projection maps, here is the one list that both built them and tests them, and here is exactly what would falsify each** (§1.2.1). Two of the five strengths — the electric charge e and the weak Fermi constant G_F — aren't independent inputs at all. They're computed from the others, so they fail against experiment if the geometry is wrong (§1.2.1). The headline is a claim you can break, not a slogan you must accept.

What earns a skeptic's hour is that the paper declares its own weakest links before any physics (§0.2), gives **exactly one** authoritative source for every load-bearing claim (§0.3), and keeps every excluded sector excluded — nonperturbative QCD, strong CP, full quantum gravity, the cosmological constant are *bounded, not used as support* (§2.2, §12). The honesty is the credibility asset: an unlabeled number, a promoted claim, a scoped-out sector smuggled back in as evidence — none of these survive the status discipline this paper inherits from the corpus.

And there is an independent reality-check that the shared geometry is *real and buildable*, owing nothing to particle physics: **try to build something with it.** The same coset whose isometries this paper turns into the four forces — $K_6 = SU(3)/T^2$, equivalently $SU(3)/U(1)^2$ — is the *admissibility chamber* of a quantum-error-correction architecture that has been designed and simulated on open-source tools (§16). A shape that is mere bookkeeping cannot be machined into a working device; a shape whose isometries *are* the gauge forces can. §16 records that bridge at its honest simulated grade — a reality-check that the geometry is buildable, not evidence that the routing is correct; it changes no force-interface status and feeds no QC number back into the physics.

Reader	Start at	Then read	What you will be able to judge
Deciding whether it's worth your hour (the skeptic)	**WHY THIS MATTERS** (above), then **§13a The Strongest Objections, Answered**	§16 (an independent reality-check that the geometry is buildable — simulated grade, changes no physics status), then §0.2–§0.3	Whether the zero-dials economy, the honesty discipline, and the engineering bridge clear your bar for a serious read
New to the routing (smart non-specialist)	The *Plain* passes of §0 and §1.2.1, then the §5.1 intuition pump, then **§0.4.1 Test It Yourself**	§5 (the four projections), then §11 (sanity targets), then §16	What is claimed, what is derived vs imported, what would kill it, and that the geometry is buildable
Physicist auditing the interface	§0.3 authority map, then §5.6 master table	§6–§9 (the four force modules) + Appendix Q certificates	Whether each projection is a real map from the active branch, at the grade stated — no more, no less
Reviewer auditing honesty	§0.2 weakest links, then §13 and **§13a**	§12 (non-claims), §6.8.9 (corpus boundary), §16's does/doesn't-license table	That every claim wears a status label, every scoped-out sector stays out, and nothing was promoted

If you have one hour and want to *break* it (hostile-review order): (1) the projection maps §5.1–§5.3 and the composite QED factorization §9.3 — if any π_i is not a real map from $\mathfrak{B}_{\text{active}}$ (or, for QED, from the electroweak EFT), the routing collapses; (2) the live electroweak anomaly witness §8.6.4 — the one quantum-consistency check you can do by hand; (3) the five-couplings provenance table §1.2.1 — verify the two derived rows (e, G_F) are routed, not fitted; (4) the scope boundary §2.2/§12 and the no-double-counting ledger §10.1 — confirm no excluded sector is used to *support* an interface claim, only to bound it. (Detail in §0.1.)

WHAT YOU CAN TAKE AWAY — IN 10 SECONDS, 5 MINUTES, AN HOUR, A WEEKEND

A reading contract: each rung is a whole, correct takeaway on its own — you are never asked to hold a promise that only pays off later.

- **~10 seconds (the thesis).** The one-line thesis, above: four forces are the symmetries of one small frozen shape, and the same shape builds a quantum computer.
- **~5 minutes (the big idea + one picture).** Read THE BIG IDEA and the §5.1 intuition pump: a *force* is a way to move the internal shape without changing it (an isometry), so the gauge groups are *counted off* a shape you are not allowed to reshape — which is why the five force-strengths carry **zero new dials** (§1.2.1).
- **~1 hour (the core chain + one check you can do by hand).** Walk §5 (the four projections) → §1.2.1 (zero-dials economy) → the **fast-check scoreboard** (§0.4, above). Then verify the one-line electroweak anomaly witness $3 \cdot \frac{1}{6} - \frac{1}{2} = 0$ (§8.6.4) yourself, and read the two worked textbook-recovery ladders — Newtonian gravity $\nabla^2\Phi = 4\pi G\rho$ (§6.6.9)

and Coulomb $E = Q/4\pi\epsilon_0 r^2$ (§9.6.3) — or run the five copy-paste checks in **§0.4.1**. At the end of the hour you can say which routings reduce to known physics, and which are honestly scoped out.

- **~a weekend (the derivations).** Read §§6–9 module by module against the §5.6 backbone-coverage table and the Appendix Q certificates, then §13/§13a for the falsifier surface. This is where the projection maps are shown to be real maps from the active branch at the grade stated — no more, no less.

Abstract

This paper routes the four forces as four projections of one frozen six-dimensional shape — the constraint-selected 13D GUT geometry developed in the companion main manuscript. The active geometric backbone is $M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1$ with $K_6 = SU(3)/T^2$; the manuscript's layer separation between base geometry, finite rule constraints, and field/operator bundles (the $\times/\oplus/\otimes$ discipline) is defined where it is first used (§4). Every gate certificate this paper relies on lives externally in Paper I (Fable_GUT); this manuscript routes, not re-derives, them (WL-3, §0.2; §14).

The purpose of this paper is narrower than a theory-of-everything claim. It shows how four familiar force sectors are routed through the same geometric backbone: gravity through the spacetime-facing metric block $G_{\mu\nu}$, the strong force through the surviving $SU(3)_c$ color action of K_6, the weak force through the S^2 spin-cover structure plus S^1_Y hypercharge, and QED/electrostatics through the unbroken $U(1)_{\text{em}}$ photon mode and its static Coulomb/Gauss-law limit. For an editor screening for scope: this is an **interface/consistency paper** — it re-closes no gate, inherits every certificate from Paper I, and is deliberately **narrower than a theory of everything** (§2.2, §12); the TOE-named sibling capsule is a separate state-ledger not under review here. Reading this paper does not require accepting any TOE claim.

The paper is an interface and consistency document. It does not replace the main GUT manuscript's gate certificates, freeze records, or review protocol, and it claims none of the sectors listed in the non-claims boundary (§2.2 — full quantum gravity, nonperturbative QCD, strong CP, cosmology, dark matter, dark energy, baryogenesis). Its claim is that the four force interfaces can be written as mutually compatible projections of one shared geometry, with explicit low-energy targets and falsification conditions. The new contribution is therefore not only the four projection maps, but the Four-Force Constraint Backbone that makes the four maps mutually auditable as one interface rather than four independent sector summaries.

Restructured-build addendum. This build organizes the paper around its own constraint list run in two tenses: the Four-Force Constraint Backbone (C-FF0–C-FF14) is both the admissibility filter that defines the four projections and the test battery the finished interfaces must keep passing (§1.2). Its headline is stated up front (§1.2.1): the five textbook force strengths — G_N, g_s, g, e, G_F — carry **zero dials of their own** in this routing, with e and G_F derived rather than independent (the full provenance is §1.2.1). Each backbone row carries a machine certificate (Appendix Q).

---

0. Reviewer First Read

> **Generalists:** skim to **§0.4** (the fast-check scoreboard) and the **§0.4.1 Test-It-Yourself** checks; §0.1–§0.3 are the hostile reviewer's entry point.

This front-matter block is the hostile reviewer's entry point. It states, before any physics, where this paper is weakest, where to attack it first, and — in one consolidated map — exactly one authoritative source for every load-bearing claim. It carries no new claim of its own; it is a routing and audit aid (added in the Phase-1.5 working copy, modeled on Paper I's Reviewer-First-Read / Known-Weakest-Links apparatus). Note on §16: the quantum-computer bridge appears because the same coset is independently buildable, not for any commercial reason; it is an EXTERNAL engineering artifact at SIMULATED grade that feeds zero numbers back into any force-interface row (C-FF13), and the physical-to-logical memory overhead it reports (~9.6–14.5:1, robust fallback 3.24:1) is an honest engineering floor, competitive with published qLDPC overheads — never an algorithm-level claim.

0.1 What to read first if you have one hour

Plain. If you are time-boxed and want to break this paper, do not start at the gravity section. Start where the interface claim is load-bearing and falsifiable on a napkin, then walk outward — the four-step hostile-review order is in the routing block above (beneath the reader-map table): projection maps §5.1–§5.3 / §9.3 → anomaly witness §8.6.4 → provenance table §1.2.1 → scope boundary §2.2/§12 + no-double-counting ledger §10.1/App. C.

Precise. The single most distinguishing structural commitment is the composite $\pi_{\mathrm{QED}} \circ \pi_{\mathrm{weak}}$ (C-FF8): if that factorization is not a real map from $\mathfrak{B}_{\mathrm{active}}$ (the frozen 13D active geometry, §3) down through the electroweak EFT, the routing collapses. Attack it first.

0.2 Known weakest links (stated before review)

The paper declares its own soft spots so a reviewer does not have to find them:

- **WL-1 — Einstein-frame conditionality (gravity).** The headline $M_{\mathrm{Pl}}^2 = M_{13}^{11}\,\mathrm{Vol}(X)$ is honest only after the volume modulus is frozen (authority external: Paper I Gate 6 / App. F). If the modulus stayed dynamical the gravity interface self-downgrades to scalar-tensor (§6.6.3.3a, §6.7.2). This is a *declared dependency on Paper I*, not a claim closed here.
- **WL-2 — Quantum self-consistency of the compactification background is inherited, not proved here.** The gravity annex's nonperturbative RG target rule (§6.8.5) and KK metric-mode quantization (§6.8.4) are model-definition level only. Whether the frozen background is quantum-stable is governed by the *corpus*, not this paper: see the TOE_FINAL Gap-04 (stabilization) discharge and the Λ verdict, cross-referenced in §6.8.9 (EXTERNAL). This paper does not lean on either result to close a four-force interface claim.
- **WL-3 — All gate certificates are external.** Every "frozen", "anomaly-free", or "certificate" statement resolves in Paper I (Fable_GUT), invoked by interface only (C-FF13). This paper re-

closes no gate; if a Paper I gate downgrades, the consuming rows here downgrade automatically (§14.3).

- **WL-4 — The strong sector's confinement/χSB/strong-CP edge is open.** $SU(3)_c$ recovery is an interface claim; confinement, the mass gap, and strong CP are Outside this paper's scope and are *not* used as support (§7.8). Strong CP in particular remains honestly underived in the corpus (§6.8.9 below; TOE_FINAL Path γ, EXTERNAL).

Fastest falsification entry point (no tools): the by-hand electroweak anomaly witness (worked in §0.4 / §8.6.4) — if the hypercharge lattice were off by one step this subtraction would fail and the charge table (§8.2) would break with it. The §0.4 scoreboard lists the rest, and §0.4.1 packages it as a copy-paste check.

0.3 Consolidated Claim-to-Authority Map (C5)

Authority rule (binding). For every load-bearing claim below, **exactly one** source is authoritative — either an in-doc location (this paper) or a named external artifact (Paper I appendix/gate, or a TOE_FINAL capsule by version label). Any other mention of the same claim elsewhere in this paper is descriptive, not load-bearing. External sources are tagged **EXT**; Paper I appendices are cited as "Paper I App. X" and resolve through that manuscript's A3.R legacy-reference ledger (§14.1). TOE_FINAL capsules are cited by stable version label (e.g. "TOE_FINAL v10"), per the campaign's cross-reference convention.

#	Load-bearing claim	Strength (this paper)	Single authoritative source	Status
A1	Four force sectors are projections of one parent geometry $M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1$	Interface claim	§5.1–§5.3 (this paper); geometry frozen at Paper I App. A1/A2 **EXT**	Interface claim
A2	Gravity is the zero-mode projection of $G_{\mu\nu}$ in Einstein frame	Interface claim (conditional on modulus freeze)	§6.3, §6.6.3.3; modulus freeze Paper I Gate 6 / App. F **EXT**	Interface claim
A3	$SU(3)_c$ survives as the left action on $K_6 = SU(3)/T^2$; QCD recovered	Interface claim	§7.3, §7.6.3–§7.6.5; reps Paper I App. D / Gates 2–3 **EXT**	Interface claim
A4	Color anomaly cancels on the shared matter content	Main-manuscript certificate	Paper I App. E/E′, Gate 5 **EXT** (witnessed in §7.6.8a)	Main-manuscript certificate
A5	$SU(2)_L$ from S^2 spin cover; $U(1)_Y$ from S^1_Y; $Q = T_3 + Y$ on every multiplet	Interface claim	§8.6.2–§8.6.3, §8.6.12 (this paper, the convention's home)	Interface claim
A6	Five electroweak anomaly classes cancel; live witness $3 \cdot \frac{1}{6} - \frac{1}{2} = 0$	Interface claim (witness) + main-manuscript certificate (full)	§8.6.4 (witness, this paper); full closure Paper I App. E/E′, Gate 5 **EXT**	Interface claim / certificate
A7	EWSB $SU(2)_L \times U(1)_Y \to U(1)_{\text{em}}$; massive W/Z, massless photon	Interface claim	§8.6.6–§8.6.8; Higgs/Wilson-line protection Paper I App. H / Gate 8 **EXT**	Interface claim
A8	Electroweak scale $v = 246.02$ GeV	Main-manuscript output (imported)	Paper I Gate 8 / App. H, Paper I §8 **EXT** (output, not re-derived here)	Main-manuscript certificate
A9	QED is the composite $\pi_{\text{QED}} \circ \pi_{\text{weak}}$; photon massless by ledger; Coulomb recovered	Interface claim	§9.3, §9.6.2–§9.6.5 (this paper)	Interface claim
A10	$e = gg'/\sqrt{g^2 + g'^2}$ and $G_F/\sqrt{2} = g^2/8M_W^2$ are derived, not independent	Interface claim (derived identity)	§9.6.7 (e), §8.6.15 (G_F) (this paper)	Interface claim
A11	Four declared numerical inputs (M_Pl; $\alpha_i^{-1}(M_Z)$; y_t; $\lvert V_{us}\rvert$); v, m_h, M_U are outputs	Main-manuscript economy (inherited)	Paper I §1.3.1 four-input table + R1.8 hashes **EXT**	Main-manuscript certificate
A12	Threshold unification target $\alpha_i^{-1}(M_Z)$; threshold vector	Main-manuscript certificate	Paper I App. G / Gate 7 **EXT**	Main-manuscript certificate
A13	Quantum self-consistency of the frozen compactification background (stabilization)	Outside this paper's scope (inherited from corpus)	TOE_FINAL v9/v10 Gap-04 discharge **EXT** (§6.8.9 below)	Outside this paper's scope
A14	Cosmological constant / vacuum-energy story	Outside this paper's scope (inherited from corpus)	TOE_FINAL v12/v14 Λ verdict **EXT** (§6.6.4, §6.8.9 below)	Outside this paper's scope
A15	Backbone coverage in both tenses (every C-FF row enforced + checked by a named Appendix Q certificate)	Interface claim	§5.6 master table + Appendix Q registry (this paper)	Interface claim

#	Load-bearing claim	Strength (this paper)	Single authoritative source	Status
A16	The active coset $K_6 = SU(3)/T^2$ (whose isometries are routed into the forces) is also an error-correction admissibility chamber, designed and simulated	EXTERNAL engineering bridge (SIMULATED grade; no interface status changes)	§16 (this paper, the bridge's home); QC corpus = Tahoe Labs design package + `SIMULATION_CAMPAIGN` / `GATE_CLOSURE` **EXT**	EXTERNAL / SIMULATED

Plain reading. The map's spine is simple: this paper owns the **routing** claims (the projection maps, the charge convention, the derived couplings, the live anomaly witness, the backbone coverage); Paper I owns every **gate certificate** (anomaly closure, thresholds, Higgs protection, the input economy, v); the **corpus capsules** (TOE_FINAL) own the two questions this paper scopes out but that have moved upstream — stabilization (Gap-04) and Λ; and the **QC engineering corpus** owns the one EXTERNAL engineering fact (A16) that the geometry is buildable, routed through §16 at simulated grade with no feedback into any interface row. No claim has two homes.

0.4 Fast checks: where the geometry gives the textbook's answer

Plain. The fastest way to decide whether this routing reduces to known physics is not to read the geometry — it is to check the **landing points**. Run each projection down to low energy and ask: does it arrive at the *same equation the textbook already has*? For three of the four sectors the answer is a worked recovery you can confirm in minutes; the table below is a routing index to those landing points, not a new claim.

A reduces-to-known-physics check must return **the same number out for the same inputs in**, and must state plainly what it does *not* claim — here, it recovers the textbook limit but does not replace GR, Maxwell, or the SM. This mirrors the posture of the companion GUT manuscript's reduces-to-known-physics certificate appendix (Paper I, *Fable_GUT*, **EXTERNAL**, §14).

Sector	Textbook target it lands on	Where it is worked	Time to check	Honest grade
Gravity	Newtonian limit $\nabla^2\Phi = 4\pi G\rho$ (weak-field, slow-motion)	§6.6.3 ladder → §6.6.9	minutes (read the ladder)	Interface claim; recovers GR's weak-field limit, does **not** replace GR
QED / electrostatics	Coulomb $E = Q/4\pi\epsilon_0 r^2$, Gauss $\nabla\cdot\mathbf{E} = +\rho/\epsilon_0$; photon massless	§9.6.2 ladder → §9.6.3, §9.6.5	minutes	Interface claim; recovers the static Coulomb/Gauss (Maxwell) limit, does **not** replace QED/Maxwell
Electroweak (consistency)	Anomaly witness $3\cdot\frac{1}{6} - \frac{1}{2} = 0$ — quantum consistency **by hand**	§8.6.4	~30 seconds, no tools	Interface claim (witness); full closure is Paper I's certificate **EXT**
Electroweak (couplings)	$e = gg'/\sqrt{g^2+g'^2}$, $G_F/\sqrt{2} = g^2/8M_W^2$ — the historical Fermi/Weinberg identities	§9.6.7 (e), §8.6.15 (G_F), §8.6.11.7 (Fermi limit)	minutes	**Derived identities of the routing, not predictions** — they *break* against measurement if the geometry is wrong (§13, C-FF4/C-FF8)

Precise. These are the paper's reduces-to-known-physics fast-checks, and the anomaly row is the sharpest: $3\cdot\frac{1}{6} - \frac{1}{2} = 0$ is one subtraction a reviewer can do on a napkin, and if the hypercharge lattice were off by a single step it would fail (§8.2 what-if). The two coupling rows are the falsifiable teeth of the zero-dials headline (§1.2.1): they are **identities**, so a reviewer who finds the standard electroweak relations is finding exactly what the routing must reproduce — not a tuned fit. What the scoreboard does *not* do is upgrade any status: gravity remains a conditional interface claim (modulus freeze, WL-1), the strong sector's confinement/strong-CP edge stays scoped out (§7.8, WL-4), and full closures live in Paper I (C-FF13). → **run these yourself: §0.4.1.**

0.4.1 Test It Yourself — Five Prompts You Can Run in Any AI

> *You don't have to trust the geometry — here are five checks you can hand to a calculator or any AI. Two are worked in this paper (run them here); three live in Paper I (Fable_GUT) and are carried as clearly tagged EXT pointers, not re-derived here. Like the §0.4 scoreboard, these are consistency/recovery checks at Interface-claim / BOUNDED grade — not proof — and they upgrade no status: they change nothing in §§5–13 (each result is already banked in the body).*
>
> *Three are Type-A (same answer with and without the geometry — equivalence/reduction); two are Type-B (a genuine Standard-Model free parameter — an INPUT, not derived — checked for consistency with data). The native pair — Prompt 2 (anomaly witness) and Prompt 3 (Newtonian + Coulomb/Poisson reduction) — is the centerpiece: both are worked in this paper.*

Prompt 2 — One-line electroweak anomaly cancellation (Type A · ≈1 min, calculator · native to this paper, §8.6.4).

> You are a particle physicist checking gauge-anomaly cancellation by hand for one Standard Model generation. Take the quarks of one generation: a color-triplet left-handed quark doublet (multiplicity 3 from color) of weak-hypercharge $Y = +1/6$, and the left-handed lepton doublet of $Y = -1/2$. For the mixed $SU(2)^2$–$U(1)$ (and gravitational–$U(1)$) anomaly, the relevant trace over the left-handed weak doublets is the sum of their hypercharges weighted by multiplicity. Compute $3 \cdot (1/6) + (-1/2)$ and state whether it vanishes. Show the one line.

Expected: $3 \cdot (1/6) - 1/2 = 0$ (i.e. $1/2 - 1/2 = 0$), exact rational arithmetic, ~30 seconds, no tools. *What it shows (Interface-claim witness; BOUNDED, by identity):* the geometry's particle list balances the ledger exactly — a consistency check, **not** a claim that the geometry is the unique origin of the SM; the cancellation is *inherited, not arranged*. This is one of **five** electroweak anomaly classes audited in §8.6.4 (the full set, incl. the mod-2 Witten condition, is there); this single line is the fastest by-hand entry point, not the whole closure. The sign is explicit so a reader cannot accidentally compute $3 \cdot (1/6) + 1/2$.

Prompt 3 — Weak-field reductions to Newton and Coulomb (Type A · ≈1 min, calculator · native to this paper — the A3 centerpiece, §6.6.9 + §9.6.3).

> You are a physicist auditing whether a higher-dimensional unification proposal reproduces textbook low-energy physics. (a) For gravity: state the weak-field, slow-motion (Newtonian) limit of general relativity as a field equation for the gravitational potential Φ sourced by mass density ρ, with

> Newton's constant G. (b) For electrostatics: state Gauss's law in differential form and the resulting Coulomb field of a point charge Q in vacuum. Write both standard results explicitly.

Expected: (a) $\nabla^2\Phi = 4\pi G\rho$ (Newtonian/weak-field limit); (b) $\nabla\cdot\mathbf{E} = +\rho/\epsilon_0$, giving $E = Q/(4\pi\epsilon_0 r^2)$. In this paper: gravity ladder §6.6.3 → §6.6.9; electrostatics §9.6.2 → §9.6.3, §9.6.5; both on the §0.4 scoreboard. *What it shows (Interface claim; recovery is BOUNDED/conditional, not a derivation):* the geometry's 4D effective Einstein–Hilbert sector and its composite QED projection *land on these same equations*; reproducing the limit is a necessary consistency check (**BOUNDED**, recorded in the body), **not** a proof of the parent geometry, and the full-GR/nonperturbative-QCD regimes are scoped out (C-FF12). *Honesty guards:* the gravity recovery is **conditional on the Paper I volume-modulus freeze** (WL-1 / §6.6.3.3a; a dynamical modulus → scalar-tensor); the Coulomb half uses the convention $\mathbf{E} = -\nabla\phi$ with $\nabla\cdot\mathbf{E} = +\rho/\epsilon_0$, so $\nabla^2\phi = -\rho/\epsilon_0$ — the minus sign sits on ϕ, not on Gauss's law (B5).

Prompt 1 — Charged-shell field energy (Type A · EXT, Paper I).

> You are a careful physicist with a calculator. A thin spherical shell carries total charge $Q = 1.00000$ C at radius $R_s = 5.00000$ m. Two concentric Gaussian surfaces sit at $r_1 = 1.00000$ m and $r_2 = 10.0000$ m. Using only standard Einstein–Maxwell electrostatics, compute the electromagnetic field energy stored in the spherical region between r_1 and r_2 (shell between the two surfaces), with $E(r) = Q/(4\pi\epsilon_0 r^2)$ for $r > R_s$ and zero for $r < R_s$, energy density $u = \epsilon_0 E^2/2$, and $\epsilon_0 = 8.8541878128 \times 10^{-12}$ F/m. Give the result to six significant figures in joules, then state the closed form.

Expected (EXT — Fable_GUT App. O.4 / O.5 Case B): $U_{\text{EM},12} = 4.49378 \times 10^8$ J, from $U = \frac{Q^2}{8\pi\epsilon_0}\left(\frac{1}{R_s} - \frac{1}{r_2}\right)$. **This paper carries only the diagnostic-only boundary (§9.8.3); the worked field-energy equivalence is EXTERNAL — Paper I App. O, routed via §0.3/§14.** *What it shows (equivalence demonstration; BOUNDED, recorded in Paper I App. O):* the geometry's $\times, \oplus, \otimes$ ledger reproduces the textbook energy object-for-object — equivalence, not proof; it does not replace GR or Maxwell.

Prompt 4 — CKM CP phase (Type B · EXT, Paper I · ≈5 min — convention-sensitive; read the −120°→+60° note first).

> You are a flavor physicist. The Standard Model does not predict the CKM CP-violating phase δ_{CKM} (equivalently the unitarity-triangle angle γ); it is a free parameter fixed only by measurement. (1) Confirm that in the SM this phase is an input, not derived. (2) Given a proposal that the phase emerges from an order-three geometric holonomy equal to $-2\pi/3 = -120°$ (raw), which after the standard Wolfenstein phase-convention alignment corresponds to a compared value of $+60.0°$, evaluate the consistency of $+60.0°$ against the measured PDG value $65.5°$ for γ. Report the discrepancy in units of the experimental standard deviation (the "pull"), assuming ~10% structural precision, and state whether $+60.0°$ is consistent with data.

Expected (EXT — Fable_GUT §7.6 table, App. I / J; scoped out of this paper, §2.2): SM cannot derive δ_{CKM}; raw holonomy $-120°$, **Wolfenstein-aligned** $+60.0°$, vs PDG $65.5°$ → **0.79σ pull** (

$(65.5 - 60.0)/6.9 \approx 0.79$). *Convention trap (read first):* the **raw** holonomy is $-120°$; only the **Wolfenstein-aligned** $+60.0°$ is compared to PDG — quoting $-120°$ against $65.5°$ is a false ~185° "mismatch." *What it shows (Type-B consistency; BOUNDED, recorded in Paper I):* sub-1σ agreement is **consistency at the stated precision, not proof**; a null would have falsified, a sub-1σ pull is a non-falsification.

Prompt 5 — Exactly three generations (Type B · EXT, Paper I).

> You are a model-building physicist. (1) Confirm that the Standard Model does not explain why there are exactly three generations of matter — the number 3 is an empirical input, not derived. (2) A proposal obtains the generation number as a topological index (a spin-ℂ / Borel–Weil–Bott index on the internal flag manifold $K_6 = SU(3)/T^2$, evaluating to -3, i.e. three chiral families), an integer that cannot be continuously tuned. Treating the family count as fixed by topology, verify that "exactly three" is consistent with experiment: state what the LEP measurement of the number of light neutrino species ($N_\nu \approx 2.984 \pm 0.008$ from the Z invisible width) and the absence of fourth-generation discoveries imply about the allowed number of standard chiral generations.

Expected (EXT — Fable_GUT §1.3.1, §3.4–§3.5, App. E): SM does not predict the generation count; geometry index $= -3 \rightarrow$ exactly three chiral families, forced (not a dial); LEP $N_\nu \approx 3$ and no fourth generation exclude the index-changing deformation that would give two or four families. *What it shows (Type-B consistency; BOUNDED, recorded in Paper I):* **consistency with the observed count, not a from-nothing proof** that nature must have three; minimality holds only inside the declared search category (a supplied natural competitor triggers the documented "category-relative diagnostic" downgrade).

---

1. Introduction

This section states the question the paper answers (§1.1 — can four force EFTs be written as compatible projections of one geometry?), the single organizing move it borrows from Paper I (§1.2 — the constraint backbone used in two tenses, as both the admissibility filter on projections and the test battery the finished interfaces keep passing), and the headline that move buys (§1.2.1 — five force strengths, one geometry, zero new dials).

1.1 The interface question

A central tension in modern unification work is that the four observed force sectors — gravity, the strong interaction, the weak interaction, and electromagnetism — are typically described by distinct effective field theories with disjoint kinematic data, disjoint symmetry groups, and disjoint conventions for matter content. Any claim that these sectors share a common origin must therefore do more than assert a parent group: it must exhibit a single geometric object together with four explicit projections, and it must show that those projections are mutually compatible at the level of state counting, anomaly cancellation, and regulator discipline.

This paper takes that posture for the constraint-selected 13D geometry $M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1$ with $K_6 = SU(3)/T^2$. The geometry, its layer separation into base manifold ($\times$), finite rule constraints ($\oplus$), and field/operator bundles ($\otimes$), and the certificate machinery that selects it are developed in the main GUT manuscript. Here we treat that backbone as given and ask a narrower question: when one writes down the projection map π_X from the active branch to each of the four force EFTs, do the four resulting interfaces fit together as views of a single object?

The answer organized in this document is yes, in a precise interface sense. The answer is controlled by a Four-Force Constraint Backbone: a compact set of interface-level constraints that require one parent geometry, shared matter and charge conventions, compatible anomaly/regulator discipline, low-energy sanity targets, and explicit falsification paths (the full fifteen rows C-FF0–C-FF14 are in §2.6). Gravity is read off the spacetime-facing metric block; the strong sector descends from the K_6 color action; the weak sector descends from the S^2 spin-cover combined with S^1 hypercharge; and QED appears as the unbroken remnant of electroweak breaking with a Coulomb/Gauss-law static limit. The same matter bundle and charge convention $Q = T_3 + Y$ feed all four projections, and the no-double-counting, anomaly, and regulator ledgers are shared.

To keep that scope honest, this paper adopts the following standalone invariant:

> A mathematically literate physicist must be able to understand the four-force routing claim from this paper alone, using only the main GUT manuscript for detailed gate certificates and frozen appendix authority.

Two consequences follow. First, the paper carries enough geometric and group-theoretic content to make the four projections precise without external lookup. Second, it does not duplicate the certificate machinery of the main GUT manuscript; whenever a load-bearing gate, anomaly cancellation, or freeze record is invoked, it is invoked by interface only, with detailed authority resting in the main GUT manuscript — **Paper I, the *Fable_GUT* build** (an earlier draft circulated as *final_manuscript*; see Section 14 for the provenance note and the full authority map). The remainder of the paper makes the scope boundary explicit (Section 2), specifies the shared geometry and layer discipline, defines the four projections, lists the cross-sector consistency conditions and their falsification map, and closes with the relation to the main manuscript.

1.2 The Inversion: The Backbone Is the Test

Plain. This paper borrows Paper I's single organizing move and instantiates it at the interface level. Section 2.6 lists the Four-Force Constraint Backbone — fifteen conditions C-FF0 through C-FF14 that any honest four-force interface must satisfy. The usual way to use such a list is as a *review checklist*: write four sector summaries, then audit them. This paper uses the same list a second way, first: as the **admissibility filter that defines the projections**. A map π_X counts as a projection of this paper only if it survives every backbone row (§5.1); the same fifteen rows then return, retrospectively, as the tests the four finished interfaces must keep passing (§13). One list, two tenses — the backbone is the test, and the test built the maps.

The obvious objection, answered before it is asked. *If the projections were constructed under the backbone, don't they pass the backbone by construction?* No, for four checkable reasons, each the interface-level analog of Paper I §1.3:

1. **Construction reads structure; the targets read nature.** The backbone's constraint-faces are structural — which parent geometry, which layer, which convention (C-FF0–C-FF8). Its evaluative faces are measured physics: the Newtonian limit, the *sign* of the QCD beta function against experiment, the Fermi constant's relation $G_F/\sqrt{2} = g^2/8M_W^2$, Coulomb's law, photon masslessness (C-FF11; Section 11). Building a map that respects layer discipline does not pre-pay a measured limit.
2. **The cross-sector meshes were never guaranteed** (each is a joint Interface claim, CONDITIONAL and separately checked, not pre-paid). Each projection could be written alone; what selection cannot arrange in advance is that all four mesh *jointly* — one matter bundle acting in four sectors (C-FF3), one charge convention surviving the weak→QED composition (C-FF4, C-FF8), and a no-double-counting ledger that closes over the *union* of all four state inventories (C-FF5; Section 10.1). These are joint conditions; any one sector's construction is blind to them.
3. **Over-determination, interface form.** The compression here is the *same-object constraint*: one frozen geometry must simultaneously source four independent EFT recoveries and five low-energy targets, with zero per-sector adjustable structure — every dial was frozen upstream in Paper I (its §4.7, §8). A dishonest "interface" buys each sector its own parent; this one is forbidden to (C-FF0), and §1.2.1 counts what that buys.
4. **The falsifiers stay live.** Every backbone row keeps a named falsification path and downgrade (Section 13, keyed row-by-row to C-FF IDs); construction cannot pass a falsifier on a projection's behalf.

The identity is again the honesty mechanism, not the loophole: one published list means there is no second standard — no construction rule hidden from the reviewer, and no evaluation rule the construction was excused from.

Precise. The flow of §2.6 makes the identity visible at both ends: $\mathfrak{B}_{\text{active}} \xrightarrow{\mathcal{C}_{\text{FF}}} \{\pi_i\} \to \{\text{EFT}_i\} \to$ sanity targets → falsification map — the list $\mathcal{C}_{\text{FF}}$ that licenses the maps at the front is the list the downgrade map evaluates at the back. The master coverage table of §5.6 assigns every row its enforcing module and its machine certificate; Section 10's three ledgers are the joint-mesh audit; Section 13 is the retrospective tense in full.

1.2.1 Five Couplings, One Geometry

Plain. Here is what the same-object constraint buys, stated once, up front. The four forces arrive in textbooks with five dial-like strengths — G_N, g_s, g, e, G_F — five numbers with five separate origin stories. In this routing they have **one origin and zero new dials**: every one is either routed from a Paper I frozen object or *computed* from the others. Two of the five are not even independent.

Coupling	Where it comes from in the routing	Provenance class
G_N	Compactification-volume relation $M_{\rm Pl}^2 = M_*^{n+2} V_K$ on the frozen geometry; $M_{\rm Pl}$ is Paper I declared input 1 (R1.8)	Routed from Paper I input 1
g_s	K_6 isometry normalization ($\mathrm{Tr}\, T^a T^b = \frac{1}{2}\delta^{ab}$) with running against the frozen threshold ledger; the measured $\alpha_3(M_Z)$ is Paper I input 2's comparison **target**, not a dial here	Routed; scored against Paper I input 2
g	S^2 spin-cover isometry, same normalization discipline; $\alpha_2(M_Z)$ likewise a target	Routed; scored against Paper I input 2
e	**Computed**: $e = gg'/\sqrt{g^2 + g'^2}$ from electroweak breaking, with g' from the $S^1_Y/\mathbb{Z}_6$ hypercharge lattice	Derived — not independent
G_F	**Computed**: $G_F/\sqrt{2} = g^2/8M_W^2$, with M_W set by $v = 246.02$ GeV — itself a Paper I *output* of the Wilson-line determinant (Gate 8)	Derived — not independent; scale is a Paper I output

Precise, with the honesty clauses. First, the provenance classes are exact: nothing in this paper introduces a coupling dial — the only measured numbers anywhere in the two-paper system remain Paper I's four declared inputs (its §1.3.1), and this table shows the four force strengths plus the Fermi scale flowing out of that same economy. Second, the two derived rows are the table's falsifiable teeth: e and G_F are *identities of the routing* — if either relation failed against measurement at the stated normalizations, the weak/QED interface would be wrong as an interface, independent of Paper I (Section 13, rows C-FF4/C-FF8). Third, the scale unification is the quiet headline: G_N's size and G_F's size are routed through one volume and one determinant living on the same frozen object — the hierarchy between gravity and the weak force becomes a statement about V_K and a winding integer, not two unrelated constants. What this paper does *not* claim is any new numerical prediction beyond Paper I's: the table is a provenance map, and its certificate is structural (Q-series, Appendix Q). The honest reading of "zero new dials" is over-determination, not a fit. A number that can break against measurement is not a tuned number.

Inputs-in / outputs-out ledger (Paper I R1.8 economy, EXTERNAL). To make the "zero new dials" headline auditable in Paper I's own hashed-table style, the table below pins what this paper's coupling rows defer to. The four numbers in the upper block are the *only* measured inputs in the two-paper system; the lower block are *outputs* of the frozen geometry that this paper imports but never re-derives. The hashes are Paper I's R1.8 freeze hashes (the authority is Paper I §1.3.1 / R0.R, EXTERNAL; reproduced here for cross-reference only).

Role	Object	Value	R1.8 hash (Paper I, EXT)	Provenance class
Input (anchor)	M_{Pl}	PDG (MEASURED)	`df5976a365c3`	sets the Planck normalization $\to G_N$
Input (comparison target)	$\alpha_i^{-1}(M_Z)$	PDG central (MEASURED)	`6a3b6ef06697`	the three couplings the threshold gate must unify $\to g_s$, g targets
Input (anchor)	$y_t(M_Z)$	$\mathbf{0.9665}$ (MEASURED → RG, PDG-derived)	`548d7099ef18`	up-sector flavor normalization (not a force-strength dial)
Input (anchor)	$\lvert V_{us}\rvert$	$\mathbf{0.22436}$ (PDG, MEASURED)	`a1bc510bc7cd`	CKM flavor anchor (not a force-strength dial)
Output (imported)	v	$\mathbf{246.02}$ GeV (Paper I output, Gate 8; imported, not re-derived here)	Paper I §8 / App. H	sets $M_W \to$ the G_F row
Output (imported)	M_Z comparison scale	$\mathbf{91.1876}$ GeV (MEASURED)	`a6852c7a6b00`	the scale at which α_i^{-1} are compared

Honesty clauses (anchors vs targets vs outputs). M_{Pl}, y_t, $\lvert V_{us}\rvert$ are **anchors** (read before comparison); $\alpha_i^{-1}(M_Z)$ is a **comparison target** (the threshold gate must hit it, it is not fitted into the geometry); v and M_Z are **outputs/inputs Paper I already froze**, imported here as compatibility constraints, never re-derived (C-FF13). The over-determination that makes this anti-fitting is Paper I's, not this paper's: four declared inputs return nineteen-plus independent frozen observables there. This paper adds **zero** further inputs — its contribution is the routing, not new numbers. These values were re-checked against the corpus during the Phase-1.5 pass and show **no drift** (the reconciliation is recorded at §6.8.9).

The worked routing for each row is its force module: G_N in Section 6, g_s in Section 7, g and G_F in Section 8, e in Section 9.

---

2. Scope and Status Boundary

This section fixes what the paper does and does not claim before any physics is done: the interface claims (§2.1), the explicit non-claims (§2.2), the four-label status vocabulary used throughout (§2.3), the global conventions ledger (§2.4), the projection-specification template every force module fills (§2.5), and the Four-Force Constraint Backbone C-FF0–C-FF14 (§2.6) that governs which projections are admissible. The scope boundary here is binding and is read as a strength, not a gap: an excluded sector is "not a failed gate" (C-FF12), and no interface claim may be closed by leaning on one.

2.1 What this paper claims

- All four force branches use the **same active parent geometry** $M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1$ with $K_6 = SU(3)/T^2$.

- Each branch has a **defined projection** $\pi_X : \mathfrak{B}_{\text{active}} \to \mathrm{EFT}_X$ from the active branch to its target effective field theory.
- The four sectors **share matter-bundle and charge conventions**, including a common $\mathcal{E}_{\text{matter}}$ assignment and the relation $Q = T_3 + Y$.
- The routing carries **no double-counting** of states, fields, or symmetry generators across sectors.
- The four projections use **compatible regulator, anomaly, and state-counting** conventions, so that thresholds, beta-function inputs, and anomaly cancellations are consistent across sectors.
- The combined result is a **unified force-interface certificate**: a single statement that the four force sectors are mutually compatible projections of one shared geometry.

2.2 What this paper does not claim

- A theory of everything (TOE).
- A full quantum-gravity UV completion.
- A full nonperturbative QCD theorem package (confinement, mass gap, hadronic spectrum, lattice equivalence).
- A solution to the strong CP problem.
- Any cosmological claim (inflation, structure formation, horizon problem).
- A dark-matter candidate or mechanism.
- A dark-energy mechanism or cosmological-constant resolution.
- A baryogenesis mechanism.
- All precision phenomenology (full CKM/PMNS fits, full electroweak precision observables, full flavor anomalies).
- The charged-shell electrostatics example as a GUT gate closure; it is illustrative of the $U(1)_{\text{em}}$ static limit, not a load-bearing gate.

2.3 Status vocabulary used in this paper

Status label	Meaning
Interface claim	The routing or compatibility statement is defined in this paper and the four sectors mesh under it.
Main-manuscript certificate	The underlying detailed gate certificate, freeze record, or frozen appendix authority lives in the main GUT manuscript and is invoked here only by interface.
Diagnostic / illustrative	The item is included for physical intuition or low-energy sanity, not as a load-bearing claim of this paper.
Outside this paper's scope	The item is not claimed here; it is either deferred to the main manuscript, to a separate companion treatment, or left open.

Earlier internal project labels (CB, GCC, SCC, WCC, ECC, rTT) and Stage 6 / Stage 7 status terminology are not used as formal statuses in this paper.

2.4 Global Conventions Ledger

The standalone-paper reading discipline requires every quantitative claim to be checkable without consulting external project files. The following ledger fixes the notational conventions used across the rest of the paper.

Convention	Paper value	Why it matters
4D spacetime notation	$M_{3,1} \equiv \mathcal{M}_4$	prevents duplicate spacetime notation
Metric signature	$(-+++)$ unless otherwise stated	fixes signs in GR/QFT formulas
Parent dimension	$13 = 4 + 6 + 2 + 1$	fixes compact-factor count
Color manifold	$K_6 = SU(3)/T^2$, $T^2 \simeq U(1)^2$	fixes quotient notation
Hypercharge convention	$Q = T_3 + Y$	fixes Y values; do not mix with $Q = T_3 + Y/2$
Matter notation	left-handed Weyl table for anomaly ledgers; Dirac notation only for low-energy QCD/QED actions	prevents rep/sign confusion
Gluon field strength	write $G^a_{\mu\nu}$, not $G_{\mu\nu}$, when color field strength is meant	avoids conflict with metric G_{MN}
Photon field	A_μ after EW breaking	distinguishes photon from generic gauge field
Gauge connection	$A^a_\mu T_a$ in $\mathcal{E}_{\text{gauge}}$	actor-layer object
Status labels	Interface claim / Main-manuscript certificate / Diagnostic / Outside scope	prevents legacy-status drift

Every force section must obey this ledger. If a section uses a different convention, it must explicitly declare the conversion.

2.5 Operational Projection Specification Template

Each force projection π_i is specified by the following data. Every force section must include a filled instance of this table.

Field	Meaning
Domain	the object being projected
Kept data	geometry / rules / fields retained
Discarded data	data intentionally ignored for this EFT
Projection operation	zero-mode average, harmonic expansion, symmetry quotient, bundle pushforward, breaking map, or static limit
Normalization	coupling / field / generator normalization
EFT target	target action or low-energy equation
Required sanity check	one observable or structural check
Falsification path	what makes the projection invalid
Paper I authority	where detailed certificate lives

Sections 6, 7, 8, and 9 each instantiate this template for π_{grav}, π_{strong}, π_{weak}, and π_{QED} respectively.

2.6 Four-Force Constraint Backbone

The four-force interface is governed by a constraint backbone inherited from the main GUT manuscript but restated here at interface level. The purpose of this backbone is to prevent the gravity, strong, weak, and QED sections from becoming four separately written effective-field-theory summaries. A force branch counts as part of this paper only if it survives the shared constraints below.

The backbone has one invariant:

> Each force sector must be recoverable as a projection of the same active branch, using the same geometry, layer discipline, matter-bundle conventions, charge normalization, state-counting rules, anomaly discipline, and status vocabulary.

These fifteen rows are this paper's one list in two tenses (§1.2): prospectively the admissibility filter on projections, retrospectively the test battery of Section 13. The constraints are not optional style rules. If any row fails, the relevant force section downgrades according to the falsification map of Section 13.

ID	Constraint	What it forces	Failure if absent
C-FF0	One-parent-geometry constraint	Every force section must descend from the same active branch $M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1$.	The paper becomes a collection of sector notes rather than a unified interface.
C-FF1	Layer-discipline constraint	Every object must be assigned to the $\times$, $\oplus$, or $\otimes$ layer.	Layer smuggling; gauge fields, finite rules, or operators are silently treated as base geometry.
C-FF2	Projection-specification constraint	Each force must fill the Operational Projection Specification table: domain, kept data, discarded data, operation, normalization, EFT target, sanity check, falsifier, and authority.	The force branch is undefined as a projection.
C-FF3	Shared matter-bundle constraint	All force sections must act on the same $\mathcal{E}_{\rm matter}$ assignments.	Different forces act on incompatible matter content.
C-FF4	Charge-normalization constraint	Weak and QED sections must use $Q = T_3 + Y$ consistently, with no hidden switch to $Q = T_3 + Y/2$.	Hypercharge/electric-charge routing becomes inconsistent.
C-FF5	No-double-counting constraint	No state, generator, field, or zero mode may be counted in two incompatible sectors.	The same mode is claimed twice or assigned two incompatible meanings.
C-FF6	Gauge-origin constraint	4D gauge connections must be treated as descended/effective bundle actors, not as primitive independent 13D bosons.	The paper contradicts the metric-first parent-geometry ontology.
C-FF7	Gravity/gauge separation constraint	The metric block $g_{\mu\nu}$, mixed blocks $G_{\mu A}$, and compact blocks G_{AB} must not be conflated.	Gravity, gauge, and internal/flavor structures become ambiguous.
C-FF8	QED-downstream-of-EW constraint	QED/electrostatics must factor through electroweak breaking; $U(1)_{\rm em}$ is not an independent parent projection.	Photon/QED is incorrectly treated as an unrelated branch.
C-FF9	Regulator/RG compatibility constraint	Running, thresholds, and beta-function statements must use shared state-counting and generator normalizations.	Cross-sector RG/unification claims become non-comparable.
C-FF10	Anomaly-compatibility constraint	Chiral matter, hypercharge, color, and weak assignments must remain compatible with the anomaly ledgers inherited from the main manuscript.	Quantum consistency of the shared matter content fails.
C-FF11	Low-energy sanity-target constraint	Each force must reproduce one recognizable 4D/low-energy target: Newtonian limit, asymptotic freedom, Fermi limit, Coulomb law, photon masslessness.	The projection lacks a recognizable physical target.
C-FF12	Scope-boundary constraint	The paper may not use excluded sectors — full QG UV completion, nonperturbative QCD, strong CP, cosmology, dark matter, dark energy, baryogenesis — to upgrade interface claims.	The companion paper overclaims beyond its interface scope.
C-FF13	Main-manuscript authority constraint	Detailed gate certificates, freeze records, and formal review vocabulary remain in the main GUT manuscript. This paper invokes them only by interface.	The companion paper pretends to re-close gates it only references.

ID	Constraint	What it forces	Failure if absent
C-FF14	Falsification-path constraint	Every projection and cross-sector integration claim must state what would make it false.	Hostile review cannot identify the first failing link.

The force-interface construction therefore follows this flow:

$$\text{active branch} \to \text{constraint backbone} \to \text{allowed projection maps} \to \text{force EFT targets} \to \text{low-energ}$$

Equivalently:

$$\mathfrak{B}_{\text{active}} \xrightarrow{\mathcal{C}_{\text{FF}}} \{\pi_{\text{grav}}, \pi_{\text{strong}}, \pi_{\text{weak}}, \pi_{\text{QED}}\} \to \{\text{EFT}_{\text{grav}}, \text{EFT}_{\text{QCD}}, \text{EFT}_{\text{EW}}, \text{EFT}_{\text{QED}}\}.$$

This explicitly makes the constraint list the backbone between the active branch and the force-specific projections.

The per-constraint enforcement map of the earlier draft is superseded by the master coverage table of §5.6, which adds the enforcing module, the machine certificate, and the falsification row for every ID.

A reviewer can therefore identify which section enforces each named constraint, and any attack on the four-force interface claim can be routed to a specific constraint ID.

---

3. The Shared 13D Geometry

This section fixes the single object all four projections read from: the parent manifold $M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1$ (§3.1), the $T^2 = U(1)^2$ notation bridge that matches Paper I (§3.2), and the statement that the geometry is taken as given here — its selection argument lives in Paper I, not in this companion (§3.3). No force section introduces a parent manifold of its own.

3.1 The parent manifold

The shared parent geometry is the 13-dimensional product manifold

$$M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1, \qquad K_6 = SU(3)/T^2.$$

In plain language, this is one 4D Lorentzian spacetime $M_{3,1}$, one 6D internal manifold K_6 that plays the role of a color/family flag manifold, one 2-sphere S^2 that carries the weak structure, and one circle S^1 that carries hypercharge with a boundary parity condition responsible for 4D chirality. All four force interfaces in this paper are projections of this single object; none introduces a separate parent manifold of its own.

3.2 The $T^2 = U(1)^2$ notation bridge

Here T^2 denotes the maximal Cartan torus of $SU(3)$. Since $T^2 \simeq U(1)^2$, $K_6 = SU(3)/T^2 = SU(3)/U(1)^2$. The T^2 notation is used throughout this paper to match the main GUT manuscript.

3.3 Why this geometry is the parent

This 13D product manifold is the active branch selected by the constraint method in the main manuscript. The selection argument — what eliminates competing branches, why the compact factorization is $K_6 \times S^2 \times S^1$ rather than some alternative, and what fixes the spin and boundary data — is not reproduced here. The present paper takes this geometry as given and shows how four familiar force EFTs arise as compatible projections of this single object.

---

4. Layer Discipline: $\times$, $\oplus$, and $\otimes$

This section states the three-layer separation every force module must respect: the three layers themselves (§4.1 — $\times$ base geometry, $\oplus$ finite chamber/admissibility, $\otimes$ field/bundle/operator content), the layered expression of the active branch (§4.2), and the binding no-smuggling rule (§4.3) that forbids relocating an object across layers. The layering is load-bearing, not bookkeeping: its necessity is the Paper I App. B2 dossier (EXTERNAL), and every object the gravity, strong, weak, and QED sections touch carries an explicit layer assignment.

4.1 The three layers

The architecture separates three distinct kinds of structure into three layers. The $\times$-layer is the base metric geometry: the product of Lorentzian spacetime with the compact internal factors, carrying the 13D metric G_{MN} and its Levi-Civita data. The $\oplus$-layer is the finite chamber / admissibility / claim-control data: discrete projectors, finite groupings, and the admissibility conditions that select which configurations count. The $\otimes$-layer is the field / bundle / Hilbert / operator structure: matter bundles, gauge bundles, Higgs sectors, and the operator algebras that act on them. The necessity of these three separate layers — that no two of them can be collapsed into one without losing content — is recorded in Paper I App. B2 (EXTERNAL; the layer-necessity dossier).

4.2 The active branch

The active branch is the layered expression

$$\mathfrak{B}_{\text{active}} = [\mathcal{M}_4 \times K_6 \times S^2 \times S^1_Y] \oplus [F^+_{\text{finite}} \oplus \mathcal{C}_{\text{admiss}}] \otimes [\mathcal{E}_{\text{matter}} \oplus \mathcal{E}_{\text{gauge}} \oplus \mathcal{E}_{\text{Higgs}} \oplus \mathcal{E}_{\text{proton}}].$$

Reading guide. The bracket $[\mathcal{M}_4 \times K_6 \times S^2 \times S^1_Y]$ is the $\times$-layer base geometry. The bracket $[F^+_{\text{finite}} \oplus \mathcal{C}_{\text{admiss}}]$ is the $\oplus$-layer finite chamber and admissibility data. The bracket $[\mathcal{E}_{\text{matter}} \oplus \mathcal{E}_{\text{gauge}} \oplus \mathcal{E}_{\text{Higgs}} \oplus \mathcal{E}_{\text{proton}}]$ is the $\otimes$-layer field content packaged as bundles over the base.

4.3 Layer-discipline binding rule

The four-force interface must respect this layer structure. No force section in this paper may smuggle a $\oplus$-layer object (a finite chamber projector or admissibility condition) or an $\otimes$-layer object (a gauge bundle, matter bundle, or Hilbert space) into the $\times$-layer geometry, and no force section may relocate a $\times$-layer object (a base-metric component) into the $\oplus$- or $\otimes$-layer. Every object referenced in the gravity, strong, weak, and QED sections must carry an explicit layer assignment, and the projection maps defined in Section 5 must respect that assignment.

---

5. Projection Maps from Geometry to Force EFTs

This section defines the four projection maps and their joint structure: the definition and admissibility condition (§5.1 — a map counts as a projection only if it survives the §2.6 backbone), the projection table and the composite QED routing (§5.2), the commuting array tying the four together (§5.3), the two bridges between the mixed metric block and the gauge bundle (§5.4–§5.5), and the backbone coverage master table (§5.6) that assigns every C-FF row its enforcing module and machine certificate. Sections 6–9 then instantiate one projection each.

5.1 Definition

Each force sector in this paper is realized as a projection from the active branch into a target effective field theory. Concretely, for each $i \in \{\mathrm{grav}, \mathrm{strong}, \mathrm{weak}, \mathrm{QED}\}$ there is a map

$$\pi_i : \mathfrak{B}_{\mathrm{active}} \to \mathrm{EFT}_i$$

that selects the layer-respecting data relevant to that force and discards the rest. The four force sections that follow (Sections 6–9) each open by naming the projection π_i for that sector and then identifying the $\times$, $\oplus$, and $\otimes$ data that survives the projection.

A map π_i is admissible in this paper only if it satisfies the Four-Force Constraint Backbone of Section 2.6. Thus a projection is not merely a notation for discarding data; it is a constrained reduction that must preserve the shared matter bundle, charge convention, layer discipline, regulator convention, and falsification path.

Intuition pump: the forces *are* the symmetries of the shape

Plain. The projection map π_i is abstract on first contact, so here is the whole mechanism in one image, with the precise statement kept beside it so the picture is an on-ramp, never a replacement. A **force**, in this picture, is not a thing added on top of the geometry — it is *a way to move the internal shape without changing it* (an isometry). Every distinct such motion is a gauge boson.

- **Smallest toy.** Take a perfect sphere. You can spin it about any axis and it looks identical afterward — the sphere has rotation symmetries. Now count them: the independent ways to spin a sphere form exactly the rotation group, and each independent spin direction is one "symmetry generator."

A *cube* has fewer (only the ninety-degree turns); a featureless blob has none. **The richer and more rigid the shape, the more symmetries it carries — and the count is fixed by the shape, not chosen.**

- **The real object.** The active internal shape here is $K_6 = SU(3)/T^2$ (§2.1, Appendix A). Its symmetries are not arbitrary: the surviving ones form exactly the $SU(3)_c$ of the strong force — **eight** independent motions, the eight gluons (§7.3). The companion S^2 factor carries the spin-cover symmetry that becomes the weak $SU(2)_L$, and the S^1_Y factor carries the single hypercharge symmetry $U(1)_Y$ (§8). Electromagnetism is then the one combination left unbroken after the weak symmetry hides itself (§9). Four forces, one shape, no force added by hand — each is read off as "a way to move $K_6 \times S^2 \times S^1$ without changing it." That reading is what the projection maps π_i make precise.
- **What is shown, and what is not.** *Shown (as an interface claim, at the grade of §§7–9):* the *gauge content* — which symmetry groups survive, how many generators each carries, the charge convention $Q = T_3 + Y$ they share — is read from the shape's isometries, not fitted. *Not shown here:* the *strengths* of those forces are a separate question (the couplings of §1.2.1, routed from Paper I anchors, not from the symmetry count), and the deep dynamics — confinement, the mass gap, full quantum gravity — stay scoped out (§2.2, §12). The honest boundary is "the forces' *symmetry structure* is geometric"; nothing past that grade is asserted.

The takeaway is: a force, here, is the answer to "how can this shape move and still be itself?" — and because a frozen shape has a *fixed* set of such motions, the gauge groups are counted off the geometry, not dialed in. This is why the paper can claim four forces from one object with **zero new dials** (§1.2.1): you cannot re-tune the symmetries of a shape you are not allowed to reshape. And if you want to see this pay off in textbook terms before reading the modules, jump to the §0.4 fast-check scoreboard: the same routing lands on $\nabla^2\Phi = 4\pi G\rho$, on Coulomb's law, and on the anomaly identity $3 \cdot \frac{1}{6} - \frac{1}{2} = 0$ — the geometry giving back exactly the equations you already trust.

5.2 Projection table

Projection	Domain	Codomain	Output
π_{grav}	$\mathfrak{B}_{\text{active}}$	gravity EFT	$g_{\mu\nu}$, stress-energy coupling
π_{strong}	$\mathfrak{B}_{\text{active}}$	QCD EFT	$SU(3)_c$, gluons, quarks
π_{weak}	$\mathfrak{B}_{\text{active}}$	electroweak EFT	$SU(2)_L \times U(1)_Y$, chiral matter
π_{QED}	electroweak EFT	QED / electrostatics	$U(1)_{\text{em}}$, photon, Maxwell/static limit

Three of the four projections take the active branch directly as their domain. The QED projection is composite: it factors through the electroweak EFT, reflecting that the unbroken $U(1)_{\text{em}}$ photon is a post-breaking remnant of $SU(2)_L \times U(1)_Y$ rather than an independent reduction of the parent geometry.

5.3 Commuting diagram

The four projections fit into a single commuting array:

$$
\begin{array}{ccccc}
\mathfrak{B}_{\text{active}} & \overset{\pi_{\text{grav}}}{\to} & \text{EFT}_{\text{grav}} & \to & \text{GR /weak field} \\
\downarrow \pi_{\text{strong}} & & \downarrow \pi_{\text{weak}} & & \downarrow \pi_{\text{QED}} \\
\text{EFT}_{SU(3)_c} & & \text{EFT}_{SU(2)_L \times U(1)_Y} & & \text{EFT}_{U(1)_{\text{em}}}
\end{array}
$$

5.4 Bridge A — Mixed metric block ↔ gauge bundle

In Kaluza–Klein block form, the 13D metric splits into a 4D spacetime block $g_{\mu\nu}$, a compact internal block g_{AB}, and a mixed block $G_{\mu A}$. The gauge zero modes of the 4D theory originate in the mixed block: components $G_{\mu A}$ associated with the compact isometries of $K_6 \times S^2 \times S^1_Y$ — equivalently, with bundle automorphisms over those factors — descend to massless 4D vector modes. In the 4D effective theory, these zero modes are packaged as gauge connections $A^a_\mu T_a$ on the principal G_{SM}-bundle $\mathcal{E}_{\text{gauge}}$. Schematically,

$$G_{\mu A} \text{ zero modes } \rightsquigarrow A^a_\mu T_a \in \mathcal{E}_{\text{gauge}}.$$

The left side is a $\times$-layer object (a component of the base metric); the right side is an $\otimes$-layer object (a connection on a bundle). The bridge says that the geometric origin of the gauge content is the mixed block, while the 4D actor that carries the gauge dynamics is the bundle connection.

5.5 Bridge B — Primitive metric vs emergent gauge actor

The 13D metric G_{MN} is the primitive bosonic object at the parent-geometry level: it is the only fundamental bosonic field in the $\times$-layer. The 4D Yang–Mills gauge fields A^a_μ are not added as independent primitive 13D fields. They are effective gauge actors obtained after compactification and after the mixed-block zero modes are repackaged, as in Bridge A, into the bundle $\mathcal{E}_{\text{gauge}}$. There is no contradiction between "the 13D metric is the only primitive boson" and "the 4D theory has Yang–Mills gauge fields" provided A^a_μ is treated consistently as descended/effective rather than primitive. The force sections that follow respect this distinction: when they speak of a 4D gauge connection, the connection lives in $\mathcal{E}_{\text{gauge}}$ at the $\otimes$-layer, not in the $\times$-layer base geometry.

5.6 Backbone Coverage Master Table

The fifteen backbone rows, each with what it demands, the module or section that enforces it under the restructured spine, the machine certificate that checks it, and its falsification row. This table is the paper's claim of completeness in both tenses (§1.2): a backbone row with no enforcing module would be an unconstructed constraint; a module discharging no row would be a sector summary smuggled past the backbone. It supersedes the enforcement sketch that closed §2.6 in the earlier draft.

ID	Demand (one line)	Enforced in	Certificate	Falsification row
C-FF0	One parent geometry for all four sectors	§3; §N.3 spec tables (domain column) of every module	Q05 (spec lint)	§13 r1
C-FF1	Every object assigned to ×, ⊕, or ⊗	§4; §N.4 routing tables; Appendix B	Q03 (ledger lint)	§13 r9
C-FF2	Full projection specification per force	§N.3 of Sections 6–9 (nine fields each)	Q05	§13 r2
C-FF3	One shared matter bundle	§8.4/§8.12 ledgers consumed by §7 and §9; §10.1	Q03	§13 r10
C-FF4	$Q = T_3 + Y$, no hidden convention switch	§8, §9 (charge tables); §2.4 ledger	Q02 (exact-fraction check)	§13 r4
C-FF5	No state, generator, field, or zero mode double-counted	§10.1; Appendix C	Q03	§13 r3
C-FF6	4D gauge fields are descended actors, never primitive 13D bosons	§5.4–5.5 bridges; §6 (metric blocks)	Q03 (home-assignment rows)	§13 r9
C-FF7	Metric block / mixed block / compact block never conflated	§6 routing and ladder	Q05	§13 r9
C-FF8	QED factors through electroweak breaking	§9.3 (composite domain); §8.8 bridge	Q04 (factorization lint)	§13 r7
C-FF9	Shared regulator / RG / state-counting conventions	§10.3; module §N.6 ladders	Q06 (targets + conventions)	§13 r5
C-FF10	Anomaly ledgers close on the shared content	§8.5/§8.13 (audit); §7.12; §10.2	Q02 (spectrum shared with Paper I G05)	§13 r6
C-FF11	One recognizable low-energy target per force	§11; worked in each module's §N.6	Q06	§13 r11
C-FF12	Excluded sectors never upgrade interface claims	§2.2; §12; module claim-boundary lines	Q07 (scope lint)	§13 r8
C-FF13	Paper I keeps certificate authority; invoked by interface only	§14 (Paper-I Interface Ledger); §N.7 authority rows	Q07	§13 r8
C-FF14	Every claim carries a stated falsifier	§13; each §N.7 closure block	Q01 (coverage lint)	§13 r12

The binding interface-module skeleton. Each force section (6–9) follows one shape, the interface analog of Paper I's constraint-module skeleton, with Section 9 (QED/electrostatics) as the calibration deep instance the others mirror at scaled depth:

```
N.1  How to read this module — ladder delta: the one or two concepts NEW to this
     force (plain + precise); everything else assumed from §§3–5 and Paper I.
N.2  What this force's EFT demands — the observed low-energy facts, plain English,
     with one concrete "what would happen if the routing were wrong."
N.3  The projection π_X — the §2.5 specification table, filled (nine fields).
N.4  What survives the projection — the ×/⊕/⊗ routing table + discarded data,
     with the no-smuggling column.
N.5  Backbone constraints carried — which C-FF rows this module discharges,
     which it consumes from other modules, per the §5.6 master table.
N.6  The math, simplified then real — a two-to-three-sentence toy, then the
     formula-by-formula ladder, ending at the module's §11 sanity target worked.
N.7  Closure — seven-line block (Claim / Mechanism / Why it matters / Paper-I
     authority / Q-certificate / Status / Falsifier), followed by the reviewer
     test suite and the module's falsification/downgrade rows.
Status labels: only the four of §2.3. A module exceeding its budget is smuggling
sector-review prose into the interface layer — cut or point to Paper I.
```

---

6. Gravity Interface

Gravity is the one force read from the parent metric directly: the spacetime-facing block of G_{MN}, with no breaking step and no bundle in between. Its module is therefore where the layer discipline earns its keep — the same primitive object feeds three channels, and the backbone rows this module carries (C-FF6, C-FF7) exist to keep those channels from being counted twice. The quantum-gravity material is collected at the end as a model-definition annex; the four-force trunk does not lean on it.

6.1 How to read this module (ladder delta)

Two concepts are new here; everything else is assumed from §§3–5 and Paper I.

One object, three channels. *Plain:* the 13D metric is the only primitive boson. Its (μ, ν) block becomes 4D gravity; its mixed (μ, A) block becomes the gauge descendants (Bridge A, §5.4); its internal (A, B) block becomes spectra, moduli, and overlap data. Nothing is added; everything is routed. *Precise:* the block decomposition of §6.6 rung 2 is the routing statement, and the no-double-counting ledger (Appendix C, rows 1–3) is its audit.

The Einstein-frame caveat. *Plain:* the headline relation $M_{\mathrm{Pl}}^2 = M_{13}^{11}\,\mathrm{Vol}(X)$ is honest only after the volume modulus is frozen (Paper I, Gate 6) and the Weyl rescaling is done. If the modulus stayed dynamical, this module would not claim Einstein gravity — it would downgrade itself to scalar-tensor, by its own table. *Precise:* §6.6 rung 3a is the conditional; §6.7.2 row 3 is the downgrade.

6.2 What gravity's EFT demands (plain English)

The low-energy facts the routing must reproduce: masses attract with an inverse-square force sourced by energy density (the Poisson equation); *everything falls the same way* — no species gets a private gravitational field (the equivalence principle, here a single stress-energy ledger on one $g_{\mu\nu}$); and the strength of the attraction, G_N, is one universal number.

One concrete what-if. The tempting wrong routing is to keep the mixed-block zero modes as gravity *and* introduce 4D Yang–Mills fields as fresh primitives. Then the same modes are counted twice (C-FF5, C-FF6 fail jointly), and worse: the gauge couplings disconnect from the compact volume, so the unification arithmetic of Paper I Gate 7 loses its geometric source. Bridge A exists to make that error unwritable — the mixed block has exactly one home.

6.3 The projection π_{grav}, specified

$$\pi_{\mathrm{grav}} : \mathfrak{B}_{\mathrm{active}} \to \mathrm{EFT}_{\mathrm{grav}}$$

This projection produces the 4D metric $g_{\mu\nu}$ and fixes the stress-energy coupling on $\mathcal{M}_4$. The projection-map definition itself is given in Section 5 of this paper; the present section unpacks the gravity-specific content of π_{grav}.

Gravity Operational Projection table.

Field	Specification
Projection	$\pi_{\mathrm{grav}} : \mathfrak{B}_{\mathrm{active}} \to \mathrm{EFT}_{\mathrm{grav}}$
Domain	$G_{MN}(x,y)$ on $\mathcal{M}_4 \times X$, $X = K_6 \times S^2 \times S^1_Y$
Kept data	zero-mode $g_{\mu\nu}(x)$, universal stress-energy coupling
Discarded data	massive KK tensor/vector/scalar modes unless explicitly retained as EFT corrections
Projection operation	normalized compact zero-mode projection plus Einstein-frame Weyl rescaling
Normalization	$M_{\mathrm{Pl}}^2 = M_{13}^{11}\mathrm{Vol}(X)$ only after volume modulus freeze/rescaling
EFT target	Einstein-Hilbert action plus declared EFT corrections
Sanity check	$\nabla^2\Phi = 4\pi G_N \rho$
Falsifier	no normalized zero mode, no Einstein-frame action, or sector-dependent leading metric coupling
Paper I authority	Paper I App. A1 (geometry), Paper I App. F (moduli stabilization), Paper I App. R0 (reproducibility) — all EXTERNAL

This projection specification instantiates the template of Section 2.5 for gravity.

6.4 What survives the projection

The gravity content of Sections 6.0–6.7 is distributed across the three layers as follows.

Layer	Gravity content	What must not be smuggled
$\times$	$\mathcal{M}_4$, 13D metric base, $G_{\mu\nu}$, compact geometry	gauge-fixing rules, Hilbert space, ghosts
$\oplus$	gauge choice, diffeomorphism quotient, BRST/BV rules, freeze/downgrade protocol	metric dimensions
$\otimes$	perturbations $h_{\mu\nu}$, stress-energy $T_{\mu\nu}$, ghost/antifield domains, physical Hilbert quotient	base geometry

The right column makes the layer separation falsifiable: an apparently valid result that secretly imports an $\oplus$-layer rule into the $\times$ layer, or base geometry into the $\otimes$ layer, is rejected even if its arithmetic is correct.

6.5 Backbone constraints carried

Backbone row	How this module carries it
C-FF7 (defining row)	The metric block $G_{\mu\nu}$, mixed blocks $G_{\mu A}$, and compact blocks G_{AB} are never conflated: §6.6 rung 2 routes them, §6.4 assigns their layers, and the correction ledger (§6.6) keeps "EFT correction" from becoming a license. Certificate: Q05.
C-FF6	4D gauge connections are descended actors via Bridge A; gravity keeps only the zero-mode spacetime metric. Certificate: Q03 (home-assignment rows).
C-FF0	Domain is $\mathfrak{B}_{\text{active}}$ directly — the first projection in the commuting array of §5.3.
C-FF1	Layer assignment in §6.4, with the no-smuggling column binding.
C-FF5	Massive KK modes are discarded-means-accounted: their counting home is the correction ledger, not a second sector. Certificate: Q03.
C-FF11	The Newtonian limit $\nabla^2\Phi = 4\pi G_N \rho$ is worked at §6.6 rung 6 (Section 11, row 1). Certificate: Q06.
C-FF12 / C-FF13	The quantum-gravity package (§6.8) is Diagnostic / model-definition only; it never upgrades an interface claim, and gate authority stays with Paper I (A1, F, R0). Certificate: Q07.
C-FF14	The eleven-row falsification map of §6.7.2, including this module's self-downgrade row (Einstein frame → scalar-tensor).

Consumes: nothing upstream — this is the trunk's first projection. **Supplies:** the shared Einstein-frame $g_{\mu\nu}$ background and the single stress-energy ledger to Sections 7–9, and the G_N row of §1.2.1 (Paper I input 1, routed through $\mathrm{Vol}(X)$).

6.6 The math, simplified then real

Toy first. Put one massless scalar on a compact space X and expand in harmonics: the integral $\int_X dy\, \sqrt{g_X}\, Y_0\, \phi(x,y)$ picks out the y-independent piece, and everything else is a tower of heavy modes. The whole gravity ladder is that one move applied to the (μ,ν) block of the metric — followed by the one honesty step the toy doesn't need: rescaling to the frame where the result *is* Einstein gravity, which is only legitimate because Paper I froze the volume modulus first.

6.6.1 Geometry-to-EFT map

Before invoking gauge-fixing, BV-BRST machinery, or RG flow, the gravity sector must specify the projection from the 13D parent metric to the 4D effective fields. The map is fixed once and reused by every subsequent subsection:

$$G_{MN}(x,y) \quad \longrightarrow \quad \begin{cases} g_{\mu\nu}(x) & \text{4D gravity} \\ A^a_\mu(x) & \text{gauge descendants from mixed blocks} \\ \phi_i(x),\, m_i,\, Y_i & \text{compact moduli / overlap / flavor data} \end{cases}$$

The projection that extracts each effective 4D field is named explicitly:

- $g_{\mu\nu}(x)$ is the normalized zero-mode projection of the (μ,ν) block against $Y_0(y) = \mathrm{Vol}(X)^{-1/2}$ on $X = K_6 \times S^2 \times S^1_Y$ (formula stated explicitly in Section 6.9.3), followed by Einstein-frame Weyl rescaling.
- $A^a_\mu(x)$ is obtained from the mixed (μ, A) block by expanding in Killing vectors of the active compact factors and retaining the zero modes associated with each gauge generator.
- $\phi_i(x)$, m_i, Y_i are obtained from the internal (A, B) block by expanding in the scalar harmonic basis of $K_6 \times S^2 \times S^1$ and recording the masses and overlap integrals that source the flavor data.

The full conversion chain enforced through Sections 6.2–6.12 is

```
13D metric -> block decomposition -> gauge quotient -> 4D effective Einstein-Hilbert sector -> stress-energy coupling ->
Newtonian / weak-field limit -> optional BRST/BV quantum interface -> claim boundary / falsification map
```

6.6.2 Classical gravity sector — interface claim

Gravity is the (μ,ν) block of the same 13D metric whose mixed block yields gauge fields and whose internal block yields spectra and overlap. Dimensional reduction of the 13D Einstein-Hilbert action over the active compact space produces 4D Einstein gravity together with the gauge kinetic sector and the scalar modulus sector. The corresponding scaling relations are

$$M_{\mathrm{Pl}}^2 = M_{13}^{11}\,\mathrm{Vol}(K_6 \times S^2 \times S^1),$$

with

$$\mathrm{Vol}(K_6) = V_0 R^6 \sqrt{x_1 x_2 x_3}, \qquad V_0 = \frac{(2\pi)^3}{\sqrt{3}}.$$

So Newton's constant and the gauge couplings descend from the same compact geometry rather than from separate input constants.

6.6.3 Formula-by-formula ladder

The seven formulas below are the operational backbone of the gravity interface. Each block is paired with a purpose statement and a "what this means" reading so the ladder is unambiguous to a reviewer who reads the section in isolation.

6.6.3.1 Parent metric

$$G_{MN}(X), \qquad M,N=0,\dots,12.$$

The purpose of this formula is to declare the 13D parent metric as the same primitive object used by every other force sector in this paper. **This means** gravity is not introduced from a separate input field; it is the spacetime-facing reading of one object that is already present in the geometry.

6.6.3.2 Block decomposition

$$G_{MN}=\begin{pmatrix} G_{\mu\nu} & G_{\mu B} \\ G_{A\nu} & G_{AB} \end{pmatrix}.$$

The purpose of this formula is to route the components of G_{MN} into the three effective channels of the EFT map. **This means** $G_{\mu\nu}$ routes gravity, $G_{\mu A}$ routes gauge descendants, and G_{AB} routes compact spectra, moduli, and overlap data, without double-counting.

6.6.3.3 Dimensional reduction

Let $Y_0(y)=\mathrm{Vol}(X)^{-1/2}$ be the normalized scalar zero mode on $X=K_6\times S^2\times S^1_Y$. The 4D metric mode is

$$g^{(0)}_{\mu\nu}(x)=\int_X d^9y\,\sqrt{g_X(y)}\,Y_0(y)\,G_{\mu\nu}(x,y).$$

The physical Einstein-frame metric $g^E_{\mu\nu}$ is obtained after the standard Weyl rescaling that fixes the coefficient of $R[g]$ to $M^2_{\mathrm{Pl}}/2$.

The purpose of this formula is to define the observed 4D metric as the normalized zero-mode projection of the spacetime-facing block, followed by Einstein-frame Weyl rescaling. **This means** the 4D gravitational degrees of freedom are not added by hand; they are the zero-mode content of $G_{\mu\nu}$ on the active compact factors, normalized against the compact measure and expressed in the frame in which $R[g]$ has the standard Einstein coefficient.

6.6.3.3a Einstein-frame caveat

The scaling relation $M^2_{\mathrm{Pl}}=M^{11}_{13}\mathrm{Vol}(X)$ is **not an unconditional prediction** — it is an Interface claim, **CONDITIONAL** on the volume modulus being FROZEN (WL-1, §0.2) or the metric transformed to Einstein frame. If the volume modulus remains dynamical, the reduction is scalar-tensor rather than pure Einstein gravity and the gravity interface must be downgraded.

6.6.3.3b Universal-coupling test

The leading reduced matter action for every matter/gauge sector must use the same Einstein-frame metric,

$$S^{(4)}_{\mathrm{matter}}=\int d^4x\,\sqrt{-g_E}\,\mathcal{L}_{\mathrm{matter}}(g^E_{\mu\nu},\Psi,A_\mu,\dots).$$

If different sectors couple to different leading metrics at leading order, the equivalence-principle interface fails.

6.6.3.4 Einstein-Hilbert target

$$S_{\rm grav}^{(4)} = \frac{1}{16\pi G_N}\int_{\mathcal{M}_4} d^4x\,\sqrt{-g}\,R[g] + S_{\rm EFT\,corr}.$$

The purpose of this formula is to state the action target that the reduced theory must hit at leading order. **This means** the reduction must recover ordinary 4D gravity to leading order, and any deviation must appear as a declared, controlled higher-order EFT correction listed in the correction ledger (§6.6.4).

6.6.3.5 Stress-energy coupling

$$G_{\mu\nu}^{(4)} = 8\pi G_N\,T_{\mu\nu}^{\rm eff}.$$

The purpose of this formula is to require all matter and gauge sectors to source the same 4D metric $g_{\mu\nu}$. **This means** there is a single stress-energy ledger on $\mathcal{M}_4$; no matter or gauge sector is allowed to source a private gravitational field.

6.6.3.6 Weak-field limit

$$g_{00} \approx -(1+2\Phi), \qquad \nabla^2\Phi = 4\pi G_N\,\rho.$$

The purpose of this formula is to demand that the reduced geometry recover Newtonian gravity in the weak-field, slow-motion limit. **This means** if the theory fails this Poisson equation it does not reproduce the low-energy regime in which gravity is already known, and the falsification / downgrade map (§6.7.2) flags the failure.

6.6.3.7 Optional BRST physical space

$$\mathcal{H}_{\rm phys} = \ker Q_{\rm BRST}/\,\mathrm{im}\,Q_{\rm BRST}.$$

The purpose of this formula is to record the physical-state definition used when the quantum interface of gravity is invoked. **This means** physical quantum states are equivalence classes under the BRST redundancy of Section 6.3; this is a companion claim and does not, on its own, close a scoped-GUT gate.

6.6.4 Correction ledger

Corrections to the leading gravity sector are tracked with explicit status and authority so that "EFT correction" is never a license to absorb arbitrary deviations.

Correction	Status	Authority
compactification scalar modes	controlled / diagnostic / excluded	Paper I App. A1 / F
higher-curvature terms	EFT correction	gravity appendix
quantum-gravity loops	outside this paper's scope	companion
cosmological constant	outside this paper's scope	Paper I App. R2 / §8; corpus state: TOE_FINAL v12/v14 Λ verdict (computed Weinberg-open), §6.8.9 below

6.6.9 The sanity target and the coupling row, worked

The C-FF11 target for this module is rung 6: $g_{00} \approx -(1+2\Phi)$ with $\nabla^2\Phi = 4\pi G_N\,\rho$ — Newtonian gravity recovered as the weak-field, slow-motion limit of the reduced Einstein equations (Section 11, row 1). And the §1.2.1 coupling row is rung 3's normalization read in reverse: $G_N^{-1} \propto M_{\rm Pl}^2 = M_{13}^{11}\,\mathrm{Vol}(K_6 \times S^2 \times S^1)$, with $\mathrm{Vol}(K_6) = V_0 R^6 \sqrt{x_1 x_2 x_3}$, $V_0 = (2\pi)^3/\sqrt{3}$ — Newton's constant is Paper I's declared input 1 ($M_{\rm Pl}$, R1.8) routed through one frozen compact volume. No gravity-sector dial exists. This is the gravity sector's reduces-to-known-physics landing point (§0.4 scoreboard): the geometry's weak-field limit is the textbook Poisson equation $\nabla^2\Phi = 4\pi G\rho$ unchanged — the same equation, not a new one — recovering GR's Newtonian limit without replacing GR.

6.7 Closure

```
Claim:      4D Einstein gravity, with universal stress-energy coupling and the
            Newtonian limit, is the normalized zero-mode projection of the
            spacetime-facing block of the one parent metric, in Einstein frame.
Mechanism:  block decomposition -> zero-mode projection against Y0 -> Weyl
            rescaling under the Paper I modulus freeze -> Einstein-Hilbert target.
Why it
matters:    anchors the routing's ontology: the metric is the only primitive
            boson, so every other force module must be descended, not added.
Paper I
authority:  App. A1/A2 (geometry), F (moduli freeze), C1/C7 (dossiers),
            R0 (reproducibility), R2 + Section 9 (claim boundary).
Certificate: Q03 (counting homes), Q05 (spec/blocks), Q06 (Newtonian target),
            Q07 (QG scope) — Appendix Q.
Status:     Interface claim (classical routing); Diagnostic (QG model-definition
            package, §6.8); Outside scope (UV completion, cosmology, Lambda).
Falsifier:  no normalized zero mode; no Einstein-frame reduction (-> scalar-
            tensor downgrade); sector-dependent leading metric; QG material
            used as gate support (§6.7.2, §13).
```

6.7.1 Minimum reviewer test suite

The following seven tests are the minimum a reviewer should be able to execute against this section. Each test is stated as a *requirement* on the gravity interface; passing the test does not upgrade any status that is already recorded, and an added test never downgrades a status already established in Sections 6.2–6.7.

Test	Required output
block decomposition test	$G_{\mu\nu}, G_{\mu A}, G_{AB}$ separated without double-counting
dimensional reduction test	explicit $g_{\mu\nu}(x)$ projection
Einstein-Hilbert target test	4D action target stated
Newtonian limit test	$\nabla^2\Phi = 4\pi G\rho$ recovered
stress-energy ledger test	all matter/gauge sectors source same $g_{\mu\nu}$
redundancy test	diffeomorphism / BRST rule stated
scope test	full QG not used as GUT gate closure

6.7.2 Falsification / downgrade map

Each gravity interface claim has a defined failure path and a defined downgrade outcome. The table is binding: a row that triggers downgrades the corresponding claim without renegotiation.

Claim	Falsification path	Downgrade
Gravity from $G_{\mu\nu}$	no defined projection to $g_{\mu\nu}$	gravity interface under-defined
Normalized metric projection	no Y_0 / no measure / no zero mode	gravity projection undefined
Einstein-frame target	no Weyl rescaling or modulus freeze	scalar-tensor only
Classical GR limit	no Einstein-Hilbert target / wrong stress-energy	classical interface fails
Gauge redundancy handled	no diffeo quotient / gauge-fixing	quantum interface invalid
Newtonian limit	no Poisson equation	low-energy gravity interface fails
Weak-field limit	fails Newtonian Poisson equation	low-energy gravity fails
Universal coupling	sector-dependent leading metric	equivalence-principle failure
Quantum-geometry row	no BRST/BV/FRG model definition	companion claim downgrades
QG material	used as GUT-gate support	scope violation
Scope boundary	full QG used to close GUT gates	scope violation

6.7.3 Authority pointers (Paper I)

Every claim in this section is anchored to a specific authority. Reviewers should treat the right column as the source of truth when a row is challenged.

Claim	Authority
$\mathcal{M}_4$	C1 / A1 (Paper I)
13D metric	A1 / A2 (Paper I)
actor domains	C7 (Paper I)
compactification / moduli	F / A1 (Paper I)
reproducibility	R0 (Paper I)
claim boundary	R2 / Section 8 (Paper I)
companion quantum-gravity material	outside this paper's scope

---

6.8 Claim boundary and the quantum-gravity model-definition annex

Plain: this annex is for the quantum-gravity specialist auditing the model-definition; the four-force trunk does not use it — a first read can skip to §7.

Everything below is collected for cross-reference; the four-force interface trunk does not require it, and none of it upgrades an interface claim (C-FF12, C-FF13).

6.8.1 Claim boundary

Binding sentence. Gravity is a $\times$-layer metric sector with $\oplus$-layer gauge redundancy rules and $\otimes$-layer physical-state / operator domains.

Claimed (interface).

- Gravity is represented by the spacetime-facing block $G_{\mu\nu}$ of the shared 13D metric.
- The 4D low-energy gravitational sector reduces to Lorentzian metric theory compatible with the Einstein-Hilbert form.
- Matter and gauge sectors couple through a shared stress-energy ledger on $\mathcal{M}_4$.
- Compactification corrections are treated as controlled EFT terms.
- Quantum-gravity artifacts, if invoked, are companion claims and not scoped-GUT gate closure.

Not claimed.

- Full quantum-gravity UV completion (**Outside this paper's scope**, OPEN).
- Cosmology, dark energy, inflation, black-hole evaporation, measurement theory.
- Proof that 13D metric theory is nonperturbatively complete — **Outside this paper's scope** (OPEN); no such proof is claimed.
- Companion quantum-gravity material does not upgrade GUT gates unless separately routed through the main manuscript's review protocol.

6.8.2 Quantum-gravity model-definition package — interface claim (model-definition level)

Diagnostic / model-definition only; not required for the four-force interface trunk.

The gravity memo fixes one explicit quantization package. This is a model-definition entry, not a proof of every quantum-gravity theorem.

Diffeomorphism / gauge redundancy. Before any gauge fixing or background split, gravity is declared modulo diffeomorphisms, $G_{MN} \sim \phi^* G_{MN}$ for $\phi \in \mathrm{Diff}(M_{13})$. Two metric configurations related by a smooth diffeomorphism of M_{13} are physically identical; the BRST/BV construction below implements this redundancy at the quantum level, and the physical state space is defined only after the quotient is taken.

Background split

$$G_{MN} = \bar{G}_{MN} + \kappa_{13} h_{MN},$$

with $\bar{G}_{MN}$ the Einstein-point compactification background.

Field content

$$\Phi = \{h_{MN}, c_M, \bar{c}_M, B_M, \Psi, \bar{\Psi}\},$$

together with antifields Φ^* in the BV formulation.

BV-BRST master action

$$S_{\mathrm{BV}}[\Phi, \Phi^*; \bar{G}] = S_{\mathrm{cl}}[\bar{G} + \kappa_{13} h, \Psi] + \int d^{13}X \, \Phi^*_A \, s\Phi^A,$$

with the master equation

$$(S_{\mathrm{BV}}, S_{\mathrm{BV}}) = 0.$$

Gauge fixing

The chosen gauge is the 13D de Donder gauge

$$F_M = \bar{\nabla}^N h_{MN} - \frac{1}{2}\bar{\nabla}_M h, \qquad h = \bar{G}^{MN} h_{MN},$$

implemented by a BRST gauge-fixing fermion of the usual form.

Canonical algebra

The 13D metric is also quantized canonically via a $12 + 1$ ADM split, with spatial metric γ_{IJ}, conjugate momentum Π^{IJ}, and equal-time commutators

$$[\hat{\gamma}^{IJ}(t, X), \hat{\Pi}^{KL}(t, X')] = i\hbar \, \delta^{(I}_K \delta^{J)}_L \, \delta^{(12)}(X - X').$$

Physical Hilbert space

The physical state space is BRST cohomology,

$$\mathcal{H}_{\rm phys} = \ker Q_{\rm BRST}/\operatorname{im} Q_{\rm BRST}.$$

6.8.3 Equivalence-principle check

The reduced matter and gauge sectors must couple universally through $g_{\mu\nu}$ at leading order. Any sector-dependent leading gravitational coupling is a failure.

Operationally, this means that the leading dimensional reduction of each matter/gauge action produces a term of the form $\int d^4x\sqrt{-g}\,\mathcal{L}_{\rm matter}[g_{\mu\nu};\Psi,A^a_\mu]$, with the same $g_{\mu\nu}$ across sectors. Sector-dependent metrics, or leading-order non-minimal couplings that distinguish matter species, are tagged as equivalence-principle failures in the correction ledger (§6.6.4).

6.8.4 Kaluza-Klein metric-mode quantization — interface claim (model-definition level)

Diagnostic / model-definition only; not required for the four-force interface trunk.

Because $K_6 \times S^2 \times S^1$ is compact, the internal mode problem is discrete. The metric fluctuation is expanded in scalar, vector, and tensor harmonics on the active compact factors, and after reduction each 4D mode coefficient obeys a canonical oscillator algebra. This is the route by which the theory defines the graviton, the vector descendants of the mixed block, the scalar modulus sector, and the higher KK excitations within one common quantization package.

6.8.5 Nonperturbative RG target rule — interface claim; nonperturbative trajectory outside this paper's scope

Diagnostic / model-definition only; not required for the four-force interface trunk.

The nonperturbative model definition uses the effective average action on the restricted theory space compatible with

$$M_{13} = M_{3,1} \times K_6 \times S^2 \times S^1, \qquad K_6 = F_2 = SU(3)/U(1)^2, \qquad \text{spin}^c(1,0) \text{ on } K_6,$$

with the selected S^2 spin-c weak sectors, the S^1 hypercharge lattice, and the active S^1/Z_2 chirality domain. The flow is governed by the exact FRG equation

$$\partial_t \Gamma_k = \frac{1}{2}\mathrm{STr}\Big[(\Gamma_k^{(2)} + R_k)^{-1}\partial_t R_k\Big].$$

The target trajectory is the one that reproduces the recovered infrared physics as $k \to 0$ and approaches a non-Gaussian UV fixed point as $k \to \infty$. Existence and detailed control of that trajectory are outside this paper's scope; the paper records only the target rule as an interface claim.

6.8.6 What the gravity section defines

Within the project's own model-definition language, the gravity interface defines the following items:

- a path integral over geometry in gauge-fixed BV-BRST form,
- a canonical commutator algebra for metric modes,
- a physical-state definition via BRST cohomology,
- a KK quantization route for the compactified spectrum,

- an explicit nonperturbative RG target rule, with fixed-point existence and global trajectory control left outside this paper's scope.

6.8.7 Gravity-sector open controls

The following items are not missing definitions; they are validation surfaces that sit outside this paper's scope. They are listed for completeness so a reviewer can see exactly which structural questions are deferred:

- constructive positivity for the interacting 13D BRST physical state space — outside this paper's scope.
- full gravitational, mixed, BV-measure, global, and S^1/Z_2 boundary anomaly closure — outside this paper's scope.
- existence and detailed structure of the required non-Gaussian fixed point — outside this paper's scope.
- global control of the RG trajectory in the restricted active theory space — outside this paper's scope.
- quantum compactification stability for $K_6 \times S^2 \times S^1$ with the active interval sector — outside this paper's scope.
- full precision global fit of the remaining geometric parameters — outside this paper's scope.

6.8.8 Scope pointer for quantum-gravity material

The quantum-gravity model-definition material in this section is collected for cross-reference but the four-force interface trunk does not require it.

6.8.9 Corpus boundary cross-reference — where the upstream questions now stand (EXTERNAL)

Diagnostic / boundary cross-reference only; not required for the four-force interface trunk. Every artifact named here is EXTERNAL to this paper.

Two structural questions that this gravity annex deliberately leaves open — (i) the **quantum self-consistency / stabilization** of the frozen compactification background, behind the nonperturbative RG target rule of §6.8.5 and the KK metric-mode quantization of §6.8.4, and (ii) the **cosmological-constant / vacuum-energy** story of the correction ledger (§6.6.4) — are governed not by this paper but by the live corpus capsule (`TOE_FINAL`). This subsection records, honestly, what changed upstream, what this paper relies on, and what stays the same in-doc. It promotes nothing: every status below is carried at the grade the corpus assigns it, by interface only (C-FF12, C-FF13).

(i) Stabilization (Gap 04) — DISCHARGED upstream; this paper still scopes it out.

Plain. The question "is the frozen internal geometry actually a stable minimum once quantum effects are included?" was open when this paper's gravity annex was first written. The corpus has since moved it forward: a stabilizing well is now reported to exist, and a Chris-countersigned ruling discharged the operative bound. This paper does **not** import that as a closed gravity claim — its gravity interface remains classical-routing-plus-model-definition (WL-2) — but the boundary is now *current* rather than silent.

Precise (EXTERNAL, by capsule label). - **TOE_FINAL v9** — *Gap 04: CLOSED at decision grade* (D15-native runner PASS under the logged countersign; package `gap04_frozen_package_native.json` **EXT**). The discharge is at **decision grade**, not full certificate grade — the corpus's own qualifier, preserved here verbatim-in-spirit. - **TOE_FINAL v10** — the boundary-sign event: with $c_{\text{bdry}} = -1.08 \times 10^{-2}$ (MEASURED-by-computation, full band) the stabilizing well is reported to **exist unconditionally** in the loop coefficient c_{loop} (the earlier c_{loop} window is vacated; self-caught correction #5). - **Dependency statement.** Forces relies on this result **only** at the level of WL-2: it is the upstream warrant that the §6.8.5 RG target rule is pointed at a background the corpus treats as stabilized-at-decision-grade. The four-force interface trunk does not consume it; no interface claim in §§6–9 is upgraded by it. What stays the same in-doc: §6.8.5 remains *model-definition level*, and §6.8.7 keeps "quantum compactification stability … outside this paper's scope".

(ii) The Λ line — Weinberg-open as a COMPUTED certificate; this paper still scopes it out.

Plain. The cosmological constant is the hardest number in physics. The corpus did not solve it — but it also no longer treats it as an open hunt: it ran the calculation that would have shown a cancellation mechanism and found none, then proved the chamber route to such a mechanism cannot work. So Λ is "open" in the specific, honest sense that a named mechanism was computed and refuted, with Weinberg's argument left standing.

Precise (EXTERNAL, by capsule label). - **TOE_FINAL v12** — the exact-weight graded supertrace over the Paper-3 §4 ledger shows **no structural cancellation at any coefficient order** (ratio 0.58 at $k = 0$, 1.000 at $k = 1$–8): the chamber grading is coefficient-blind. **Λ stays Weinberg-open** as a *computed one-loop certificate*, not an open mechanism hunt. - **TOE_FINAL v14** — *I3 THEOREM_REFUTED*: the Λ operator is the unit operator (grading-even, label-blind); the chamber route to Λ is **closed**; Weinberg stands. No observed-Λ comparison was performed. - **Dependency statement.** Forces relies on this result **only** to keep its own boundary honest: the correction ledger's "cosmological constant — outside this paper's scope" row (§6.6.4) now names the *current* corpus state of that question rather than deferring vaguely. Nothing in this paper implies a live Λ-cancellation mechanism; the corpus has closed that route as an honest fail. What stays the same in-doc: Λ / cosmology remain Outside this paper's scope (§2.2, §12), unchanged.

Strong CP (boundary note, for completeness). The strong-sector annex scopes out strong CP (§7.8.3). The corpus likewise still has $\bar{\theta}$-smallness **honestly underived** (TOE_FINAL Path γ, "closed-as-read … θ̄-smallness honestly underived → nEDM probes it", **EXT**). The boundary is therefore current and unchanged: Outside this paper's scope, not a closed claim.

Sanity-target consistency check (MR-2 / coupling rows). The frozen four-input headline that this paper's coupling rows defer to (Paper I §1.3.1 **EXT**: M_Pl; $\alpha_i^{-1}(M_Z)$; $y_t(M_Z) = 0.9665$; $|V_{us}| = 0.22436$) and the imported electroweak scale $v = 246.02$ GeV (Paper I Gate 8 output **EXT**) were re-checked against the corpus during this Phase-1.5 pass: **no drift** was found — the §1.2.1 provenance table, the §8.5 import, and the §11 sanity targets agree with the GUT four-anchor headline as cited. No number in this paper was changed; this line records that the reconciliation was performed (C5 authority A8/A11/A12).

One-line summary (binding). Upstream, stabilization is discharged-at-decision-grade and Λ is a computed Weinberg-open certificate; in this paper, both remain **Outside this paper's scope**, now cited to their governing capsule. Promotions in this subsection: zero.

7. Strong / QCD Interface

The strong module is where "force from geometry" is most literal: the color symmetry is not assigned to K_6 — it *survives* the construction of K_6, as the left action remaining on $SU(3)/T^2$. Its honesty hinge is notational: the action is written with Dirac quarks, the anomaly ledger with left-handed Weyl fields, and the module's binding rule is that these are two notations for **one** spectrum.

7.1 How to read this module (ladder delta)

Two concepts are new here; everything else is assumed from §§3–5 and Paper I.

Color from a quotient that keeps its symmetry. *Plain:* dividing $SU(3)$ by its maximal torus removes two phase directions but does not remove the symmetry — the full $SU(3)$ still acts from the left on what remains. That surviving action is what descends to the eight gluons. *Precise:* $g \cdot [h] = [gh]$ on $K_6 = SU(3)/T^2$, $\dim = 8 - 2 = 6$ (§7.6 rungs 1–3); the gauge field is the adjoint descendant of this isometry algebra via Bridge A.

Two notations, one spectrum. *Plain:* low-energy QCD reads most naturally with Dirac quarks q_f; anomaly accounting is only honest with left-handed Weyl fields ($u_R \to u_L^c$ in $\bar{\mathbf{3}}$). Converting between them is bookkeeping — but doing the anomaly sum in the wrong notation is the classic way to fake a cancellation. *Precise:* §7.6's representation ledger and its 7.12a Weyl companion are the same content; the per-generation $SU(3)^3$ cancellation ($2 \times \mathbf{3}$ from Q_L against $u_L^c \oplus d_L^c$) must be computed in the Weyl convention (C-FF10).

7.2 What QCD demands of the routing (plain English)

The low-energy facts: eight self-interacting gluons (light-on-light is absent in QED; gluon-on-gluon is the defining feature here); quarks in color triplets, three families; a coupling that *weakens* at short distance — asymptotic freedom, the experimentally settled sign; and exact color anomaly cancellation, without which the quantum theory is simply inconsistent.

One concrete what-if. Suppose the routing assigned right-handed quarks plain triplets instead of conjugate triplets in the Weyl ledger. The Dirac-notation action would look identical — and the $SU(3)^3$ anomaly would silently fail to cancel. Nothing visible breaks at tree level; the theory dies at one loop. That is why this module's anomaly rows defer to a frozen, convention-locked ledger (Paper I Gate 5 / App. E′, EXTERNAL) rather than to inspection of the action.

7.3 The projection π_{strong}, specified

$$\pi_{\text{strong}} : \mathfrak{B}_{\text{active}} \to \mathrm{EFT}_{SU(3)_c}$$

Field	Specification
Projection	$\pi_{\text{strong}} : \mathfrak{B}_{\text{active}} \to \mathbf{EFT}_{SU(3)_c}$
Domain	$K_6 = SU(3)/T^2$, $\mathcal{E}_{\text{gauge}}$, $\mathcal{E}_{\text{matter}}$
Kept data	left $SU(3)$ action, color gauge connection, quark color modules
Discarded data	weak / hypercharge labels except where needed to identify SM multiplets
Projection operation	retain K_6-sourced $SU(3)$ zero modes and color representation modules
Normalization	$\mathrm{Tr}(T^a T^b) = \frac{1}{2}\delta^{ab}$, unless main manuscript declares otherwise
EFT target	QCD Yang-Mills + quark action
Sanity check	$\beta(g_s) < 0$ for SM active flavors
Falsifier	wrong color reps, nonzero color anomaly, or missing non-Abelian field strength
Paper I authority	Paper I App. C2 (K_6), Paper I App. D (SM recovery), Paper I App. C7 / E (matter & anomaly) — all EXTERNAL

This projection specification instantiates the template of Section 2.5 for the strong sector.

This projection extracts the unbroken $SU(3)_c$ color sector from the shared 13D geometry and yields the 4D QCD action together with its representation content. The projection-map definition itself is given in Section 5 of this paper; the present section unpacks the strong-specific content of π_{strong}.

7.4 What survives the projection

The strong-sector content distributes across the three architectural layers as follows:

Layer	Strong-sector content
$\times$	$K_6 = SU(3)/T^2$, compact spectrum, color geometry
$\oplus$	center/charge rules, regulator, freeze discipline, color-singlet admissibility, θ-scope rule
$\otimes$	quark fields, color modules $V_{SU(3)}$, gluon connection, field strength, BRST/ghost domains

7.5 Backbone constraints carried

Backbone row	How this module carries it
C-FF0	Domain is $\mathfrak{B}_{\text{active}}$ directly; color descends from the K_6 left action — weak or hypercharge generators are never reused as color.
C-FF10 (defining row)	The representation/anomaly ledgers of §7.6, in the locked Weyl convention, must close on the shared matter content; authority is Paper I Gate 5 / E′. Certificate: Q02 (spectrum shared with Paper I G05).
C-FF3	Quarks carry color through the same $\mathcal{E}_{\text{matter}}$ the weak module charges — one bundle, two sectors (no private quark copies). Certificate: Q03.
C-FF1	Layer assignment in §7.4; the θ-scope rule and singlet admissibility live at ⊕, never smuggled into ×.
C-FF5	Quark modes are counted once: physical matter in the action *and* anomaly-trace contributors is a permitted dual role under the state-counting rule (Appendix C); two physical copies are not. Certificate: Q03.
C-FF9	The β-function statement uses the shared state-counting and $\mathrm{Tr}(T^aT^b) = \frac{1}{2}\delta^{ab}$ normalization (§2.4); thresholds defer to Paper I G.x ledgers.
C-FF11	Sanity target: $\beta(g_s) < 0$ for SM active flavors, worked at §7.6 rung 7 (Section 11, row 2). Certificate: Q06.
C-FF12 / C-FF14	Confinement, mass gap, χSB, strong CP: Outside scope by the §7.8 scope box, with falsification rows live in §7.7.2. Certificate: Q07.

Consumes: the chiral matter table (the quark sector's 4D chirality uses the $S^1/\mathbb{Z}_2$ no-mirror lift — Paper I Gate 4 — not K_6 alone) and the shared bundle from Section 8's ledgers. **Supplies:** color modules to the joint anomaly audit (§10.2) and the g_s row of §1.2.1 (normalization fixed here; $\alpha_3(M_Z)$ is Paper I's comparison target through the frozen threshold ledger, Gate 7).

7.6 The math, simplified then real

Toy first. Take the circle as a baby quotient: $U(1)$ acting on itself leaves nothing, but $SU(2)/U(1) = S^2$ keeps a full $SU(2)$ worth of rotations acting on the sphere. The color story is the same move one rank up: quotient $SU(3)$ by its torus and the *whole* group still acts — six dimensions of carrier, eight surviving symmetry directions. The ladder below is that observation made into a gauge sector.

7.6.1 Geometry-to-QCD map

The strong-sector pipeline runs through the following named stages, each of which is referenced by the explicit formulas in Section 7.11 and the test rows in Section 7.14:

```
K_6 = SU(3)/T^2
  -> left SU(3) action
  -> Lie algebra su(3)
  -> adjoint gauge field A^a_mu
  -> field strength G^a_munu
  -> quark matter bundle with V_SU(3)
  -> 4D QCD action
  -> RG / anomaly / color-singlet checks
```

Every later subsection lives on one of these arrows: §7.3 supplies the action and algebra step, §7.4 supplies the matter-bundle step, §7.5 supplies the 4D QCD action step, §7.6 supplies the RG/asymptotic-freedom step, §7.12 supplies the anomaly-trace step, and §7.7 plus §7.8 sit on the optional nonperturbative companion side of the pipeline rather than on the interface-claim trunk.

7.6.2 Strong-sector claim

The strong interface claim is that the low-energy color EFT obtained from the active branch has the standard $SU(3)_c$ Yang-Mills form, with quark representation content matching the main-manuscript matter ledger. In the unified view, the strong force is the color projection of the same geometric transport law that already produces gravity and electroweak structure.

7.6.3 Color from internal isometries — interface claim

Because the compact manifold is the full flag manifold with built-in $SU(3)$ structure, the mixed metric block supports an unbroken adjoint color sector in 4D. The eight color directions are identified with the adjoint descendant of the internal isometry algebra.

Color charge is then geometric: it is internal momentum or weight data along the corresponding Killing directions.

7.6.4 Quark matter and three families — interface claim for QCD recovery; nonperturbative QCD outside scope

The low-energy quark sector has three factor-level family modes from the K_6 spin-c index. Its 4D chirality is not a K_6-only result; it uses the selected five-bundle table together with the active S^1/Z_2 no-mirror lift. The strong-sector note then freezes the low-energy quark assignment so that the reduced 4D chiral content matches the Standard Model color structure generation by generation.

Within the unified theory, the same quark fields that participate in the electroweak embedding also furnish the color triplet matter content of the strong sector. The strong-sector anomaly ledger balances generation by generation once the low-energy chiral table is fixed.

7.6.5 4D QCD descendant — interface claim for QCD recovery

After compactification and low-energy reduction, the strong sector takes the standard QCD form,

$$S_{\mathrm{QCD}} = \int d^4x \sqrt{-g_4} \left[-\frac{1}{4} G^a_{\mu\nu} G^{a\mu\nu} + \sum_f \bar{q}_f (i\gamma^\mu D_\mu - m_f) q_f \right],$$

where the masses m_f are not primitive inputs but descendants of the overlap geometry.

7.6.6 Asymptotic freedom — interface claim for QCD recovery

The strong-sector note makes the asymptotic-freedom sign explicit. Given the rep content recovered in §7.6.3–§7.6.5 — exact non-Abelian color with the usual chiral quark content — the one-loop coefficient carries the QCD sign and the coupling weakens at high energy. This is a necessary consistency check the routing passes, not an independent prediction of asymptotic freedom: the recovered color sector reproduces the standard sign because the representation content was fixed upstream (§7.6.3–§7.6.5), and the one-loop coefficient is a function of that content alone.

QCD beta-function sanity check. Explicitly,

$$\beta(g_s) \;=\; -\frac{g_s^3}{16\pi^2}\Big(11 - \tfrac{2}{3}\,n_f\Big) + \cdots,$$

which is asymptotically free for the SM number of active quark flavors.

7.6.7 Formula-by-formula ladder

The seven formulas below are the load-bearing identities of the strong-sector pipeline. Each is wrapped with its purpose and its physical meaning so that no reviewer has to reverse-engineer the role of any single expression.

7.6.7.1 Color geometry

$$K_6 \;=\; SU(3)/T^2.$$

The purpose of this formula is to identify the six-dimensional compact carrier that hosts the color branch. **This means** the color branch is routed through the six-dimensional flag manifold of $SU(3)$, not through an ad hoc internal space.

7.6.7.2 Dimension and quotient

$$\dim_{\mathbb{R}}\big(SU(3)/T^2\big) \;=\; 8-2 \;=\; 6.$$

The purpose of this formula is to check the surviving real dimension after quotienting by the maximal torus. **This means** six color/family directions remain after the two torus phase directions are removed; this is the dimension count K_6 needs in order to play the compact-internal role.

7.6.7.3 Surviving action

$$g\cdot[h] \;=\; [gh], \qquad g,h\in SU(3).$$

The purpose of this formula is to show that the quotient still supports a left $SU(3)$ action. **This means** color symmetry is not killed by the quotient; the full $SU(3)$ survives as a left action on K_6, which is what later produces the gauge symmetry.

7.6.7.4 Gauge connection

$$A_\mu \;=\; A_\mu^a T^a, \qquad T^a\in\mathfrak{su}(3).$$

The purpose of this formula is to define the 4D gluon field as a Lie-algebra-valued connection. **This means** the 4D gluon field takes values in the Lie algebra of the surviving color symmetry, with eight generators T^a matching the eight gluons.

7.6.7.5 Field strength

$$G^a_{\mu\nu} \;=\; \partial_\mu A^a_\nu - \partial_\nu A^a_\mu + g_s f^{abc} A^b_\mu A^c_\nu.$$

The purpose of this formula is to record the non-Abelian field strength of the surviving color connection. **This means** gluons self-interact through the commutator term, which is the signature of a genuinely non-Abelian gauge sector rather than a copy of QED.

7.6.7.6 QCD action

$$S_{\rm QCD} \;=\; \int d^4x\,\sqrt{-g_4}\,\left[-\tfrac{1}{4}G^a_{\mu\nu}G^{a\mu\nu} + \sum_f \bar q_f(i\gamma^\mu D_\mu - m_f)q_f\right].$$

The purpose of this formula is to state the recovered 4D action for the strong sector. **This means** standard QCD is reproduced, with quark masses m_f inherited from the flavor/overlap geometry rather than inserted as primitives.

7.6.7.7 Beta-function sanity check

$$\beta(g_s) \;=\; -\frac{g_s^3}{16\pi^2}\left(11 - \tfrac{2}{3}\,n_f\right) + \cdots.$$

The purpose of this formula is to check the RG sign of the recovered coupling. **This means** the recovered theory is asymptotically free for the SM number of active quark flavors, matching the standard QCD running.

7.6.8 Representation and anomaly ledger

Field	$SU(3)_c$ rep	Chiral?	Anomaly contribution	Authority
Q_L	$\mathbf{3}$	yes	included	D/E
u_R	$\mathbf{3}$ or conjugate convention	yes	included	D/E
d_R	$\mathbf{3}$ or conjugate convention	yes	included	D/E
leptons	singlet	yes	none for color	D/E

The strong-sector claim fails if representation conventions make color anomalies nonzero or mismatch the matter bundle.

7.6.8a Left-handed Weyl representation ledger

Left-handed Weyl field	$SU(3)_c$ rep	Role
Q_L	$\mathbf{3}$	left quark doublet
u^c_L	$\bar{\mathbf{3}}$	conjugate of right up quark
d^c_L	$\bar{\mathbf{3}}$	conjugate of right down quark
L_L	$\mathbf{1}$	lepton doublet
e^c_L	$\mathbf{1}$	conjugate of right electron

The low-energy QCD action may use Dirac quarks q_f, but the anomaly ledger must use left-handed Weyl fields. These are notation changes, not spectrum changes.

Per generation, the apparent $SU(3)^3$ anomaly contribution from the two color-triplet components of Q_L is cancelled by u_L^c and d_L^c. Color anomaly cancellation must be checked in the same representation convention used by the matter ledger.

7.6.9 Strong-sector result

The strong force is therefore represented as:

- unbroken $SU(3)_c$ descending from the same mixed metric block,
- geometric color charge,
- quark matter inherited from the three factor-level family modes plus the active no-mirror chirality lift,
- QCD as the reduced 4D descendant,
- asymptotic freedom in the ultraviolet,
- color-singlet representation channels and chiral-breaking posture in the infrared, with the nonperturbative confinement / mass-gap and chiral-dynamics structure outside this paper's scope.

7.6.10 The sanity target and the coupling row, worked

The C-FF11 target is rung 7: $\beta(g_s) = -\frac{g_s^3}{16\pi^2}\left(11 - \frac{2}{3}n_f\right) + \cdots < 0$ for the SM active flavor count — asymptotic freedom with the measured sign, as a consequence of the recovered non-Abelian content rather than an extra assumption (Section 11, row 2). The §1.2.1 coupling row: g_s carries no dial of its own — its generator normalization is the §2.4 convention, its running uses the shared state counting, and its measured value $\alpha_3(M_Z)$ enters only as Paper I's declared comparison **target** (R1.8, input 2), scored through the frozen threshold ledger of Paper I Gate 7.

7.7 Closure

```
Claim:      the low-energy color sector of the active branch is standard SU(3)_c
            QCD: surviving left action -> adjoint gluons -> triplet quarks on the
            shared bundle -> QCD action -> asymptotic freedom, anomaly-free.
Mechanism:  K6 = SU(3)/T^2 isometry descent (Bridge A) + shared-bundle matter
            + locked Weyl anomaly ledger (Rungs 1-7).
Why it
matters:    the non-Abelian half of the routing claim; shows the same geometry
            that yields an Abelian remnant (Sec. 9) yields a confiningly-signed
            non-Abelian sector with the right representation arithmetic.
Paper I
authority:  App. C2/A1 (K6), D/C7/A2 (color reps), E' (anomaly), G (thresholds),
            I/J (quark masses as descendants), R0 (reproducibility).
Certificate: Q02 (reps/anomaly on shared spectrum), Q03 (counting homes),
            Q06 (beta-sign target), Q07 (scope) — Appendix Q.
Status:     Interface claim (QCD recovery); Main-manuscript certificate (anomaly,
            thresholds, masses); Outside scope (confinement, mass gap, chiSB,
            strong CP — §7.8).
Falsifier:  no surviving SU(3) action; wrong color reps; nonzero color anomaly;
            beta-function sign wrong; or any excluded item used as support
            (§7.7.2, §13).
```

7.7.1 Minimum reviewer test suite

Test	Required output
quotient test	$SU(3)/T^2$ defined and six-dimensional
action test	surviving $SU(3)$ action shown
connection test	$A_\mu^a T^a$ defined
representation test	all quarks have correct color reps
anomaly test	color/mixed anomalies vanish
beta-function test	asymptotic freedom sign correct
scope test	confinement/strong CP not overclaimed

7.7.2 Falsification / downgrade map

Claim	Falsification path	Downgrade
K_6 sources color	no surviving $SU(3)$ action	Gate 2 fails
correct quark color reps	wrong reps or missing matter modules	D/E fail
QCD action recovered	no non-Abelian field strength / quark coupling	strong interface fails
anomaly cancellation	nonzero $SU(3)$ anomaly	Gate 5 fails
asymptotic freedom	beta function sign wrong	QCD consistency fails
confinement / mass gap	not routed externally	diagnostic only
strong CP	not routed externally	outside this paper's scope

7.7.3 Authority pointers (Paper I)

Claim	Authority
K_6	C2 / A1 (Paper I)
color reps	D / C7 / A2 (Paper I)
anomaly	E (Paper I)
threshold	G (Paper I)
flavor masses	I/J (Paper I)
reproducibility	R0 (Paper I)
companion QCD nonperturbative material	outside this paper's scope

---

7.8 Claim boundary and nonperturbative scope annex

Everything below bounds the module; none of it is load-bearing for the interface trunk, and the scope box's statuses are binding (C-FF12).

7.8.1 Claim boundary

The strong-sector content of this section makes a deliberately bounded set of interface claims and explicitly defers a second set as outside this paper's scope.

Interface claims for QCD recovery:

- $K_6 = SU(3)/T^2$ is the compact color/family carrier (equivalently $K_6 = SU(3)/U(1)^2$, since $T^2 = U(1)^2$; the existing prose elsewhere in this paper uses the $SU(3)/U(1)^2$ form).
- $SU(3)$ survives the quotient as a left action.
- The 4D descendant contains an $SU(3)_c$ gauge connection.
- Quarks carry color representations via the shared matter bundle.
- The QCD kinetic and quark coupling terms are recovered.
- Asymptotic freedom is consistent with the recovered matter content.
- Color anomalies vanish as part of the anomaly ledger.

Outside this paper's scope:

- Confinement proof — **Outside this paper's scope** (OPEN); none is claimed.
- QCD mass gap.
- Full hadron spectrum.
- Chiral symmetry breaking.
- Strong CP solution.

This paper does not use confinement, the QCD mass gap, hadron spectroscopy, chiral symmetry breaking, or strong CP to establish the strong interface. Those topics remain outside the present paper. They may be treated elsewhere, but they are not companion requirements for this four-force routing claim.

7.8.2 Confinement, singlets, and hadrons — outside this paper's scope

At low energy, confinement and color-singlet organization are required for the observed hadronic spectrum. A controlled nonperturbative route for this sector would supply a center-sector / block-spin construction, a Wilson-loop area law, a positive string tension, a mass gap, a controlled singlet-spectrum route, continuum/volume diagnostics, and a chiral-dynamics audit. Observable asymptotic states would then be treated with the standard color-singlet posture.

This paper does not use confinement, the QCD mass gap, hadron spectroscopy, chiral symmetry breaking, or strong CP to establish the strong interface. Those topics remain outside the present paper. They may be treated elsewhere, but they are not companion requirements for this four-force routing claim.

7.8.3 Strong CP note — outside this paper's scope

The strong-CP problem — i.e., why the effective $\bar{\theta}$ parameter is constrained to be extremely small — is outside this paper's scope. An APS / Z_2 counterterm calculation cancelling the quark determinant phase under a specified topological domain would be one route to address this question elsewhere; this paper does not use confinement, the QCD mass gap, hadron spectroscopy, chiral symmetry breaking, or strong CP to establish the strong interface. Those topics remain outside the present paper. They may be treated elsewhere, but they are not companion requirements for this four-force routing claim.

7.8.4 Strong CP and confinement scope box

> The recovered $SU(3)_c$ branch is sufficient for Standard Model color recovery and the QCD interface. Nonperturbative confinement, the QCD mass gap, hadron spectroscopy, and the strong-CP problem are not used to close the scoped GUT gates and are outside this paper's scope.

Companion claim	Status (this paper)	Required artifact (if independently routed)	If unrouted
confinement	outside this paper's scope (OPEN)	named nonperturbative run/proof (OPEN)	diagnostic only
mass gap	outside this paper's scope (OPEN)	named nonperturbative run/proof (OPEN)	diagnostic only
strong CP	outside this paper's scope	APS / Z_2 counterterm audit	outside this paper's scope
chiral symmetry breaking	outside this paper's scope	controlled nonperturbative calculation	diagnostic only

The "outside this paper's scope" status above is the routing label used for the interface trunk. This paper does not use confinement, the QCD mass gap, hadron spectroscopy, chiral symmetry breaking, or strong CP to establish the strong interface. Those topics remain outside the present paper. They may be treated elsewhere, but they are not companion requirements for this four-force routing claim.

8. Weak / Electroweak Interface

The weak module carries the heaviest interface load in the paper: it owns the shared matter table that Sections 7 and 9 act on, it executes the breaking that Section 9's whole existence depends on, and it hosts the one *live arithmetic witness* in the routing — the per-generation anomaly sum a reviewer can check on a napkin. Two of the five coupling rows of §1.2.1 are worked here.

8.1 How to read this module (ladder delta)

Two concepts are new here; everything else is assumed from §§3–5 and Paper I.

The spin cover — why a sphere yields doublets. *Plain:* the rotations of S^2 form $SO(3)$, which has no doublets. Weak doublets live on the double cover: $SU(2) \to SO(3)$. The sphere is a true weak-isospin carrier only through its spin-c structure — that one cover step is what turns geometry into $SU(2)_L$ representations. *Precise:* §8.6 rung 2; the doublet content then comes from the S^2 spin-c monopole sectors (Paper I, dossier C3).

Chirality as boundary data. *Plain:* handedness is not assigned field-by-field; it is enforced by where fields can live. The hypercharge circle carries a parity identification $\theta \sim -\theta$, and the $\mathbb{Z}_2$ boundary data kills the would-be mirror partner of every chiral mode. Parity violation is structural — the geometry has no place for the wrong-handed states. *Precise:* §8.6 rung 3 plus the no-mirror lift (Paper I Gate 4 / App. E, EXTERNAL); mirror survival is a named falsifier in §8.7.2.

8.2 What the weak EFT demands (plain English)

The low-energy facts: the weak force violates parity maximally (only left-handed particles feel the charged current — the 1957 shock, here a routing consequence); its carriers $W^\pm, Z$ are *massive* while the photon stays massless, so one breaking step must split one multiplet's fate; every multiplet's electric charge comes out of one formula, $Q = T_3 + Y$; at low energy the whole sector collapses to Fermi's four-fermion interaction with one constant G_F; and the quantum theory must cancel five separate anomaly classes on the same matter content — no spare fields allowed.

One concrete what-if. Shift one hypercharge by one lattice step — say $Y(L_L) = -1/3$ instead of $-1/2$. Two things fail at once: the charge table gives the electron $Q = -5/6$, and the napkin witness breaks, $3 \cdot \frac{1}{6} - \frac{1}{3} \neq 0$, so the $SU(2)^2U(1)$ anomaly is nonzero. The same number that fixes the charges fixes the quantum consistency: that double duty is why the hypercharge lattice has no slack (C-FF4, C-FF10).

8.3 The projection π_{weak}, specified

$$\pi_{\text{weak}} : \mathfrak{B}_{\text{active}} \to \mathrm{EFT}_{SU(2)_L \times U(1)_Y}.$$

Weak Operational Projection table.

Field	Specification
Projection	$\pi_{\text{weak}} : \mathfrak{B}_{\text{active}} \to \mathbf{EFT}_{SU(2)_L \times U(1)_Y}$
Domain	S^2, $S^1_Y/\mathbb{Z}_2$, $\mathcal{E}_{\text{matter}}$, $\mathcal{E}_{\text{Higgs}}$
Kept data	weak doublets/singlets, hypercharge, Higgs/Wilson-line doublet
Discarded data	color except for quark identity
Projection operation	spin-cover projection + hypercharge-line projection + chiral boundary projection
Normalization	$Q = T_3 + Y$
EFT target	electroweak EFT
Sanity check	charge table, anomaly cancellation, $m_W/m_Z/\rho$, Fermi limit
Falsifier	wrong charge, mirror survival, anomaly, no $U(1)_{\text{em}}$ remnant
Paper I authority	Paper I App. C3 (S^2), Paper I App. C4 ($S^1_Y/\mathbb{Z}_2$), Paper I App. D (SM recovery), Paper I App. H (Higgs protection), Paper I App. K (lepton/neutrino) — all EXTERNAL

This projection specification instantiates the template of Section 2.5 for the weak sector.

8.4 What survives the projection

The weak-sector content is routed across the three architectural layers as follows.

Layer	Weak-sector content
×	S^2, $S^1_Y/\mathbb{Z}_2$, compact spectra
⊕	chirality/parity rules, $\mathbb{Z}_2$, $\mathbb{Z}_6$, projectors, $Q = T_3 + Y$, freeze protocol
⊗	weak gauge fields, hypercharge field, matter doublets/singlets, Higgs/Wilson-line operator, lepton/quark domains

The × layer is the compact-geometry carrier set. The ⊕ layer is the algebraic/projector data that selects chirality and the charge formula. The ⊗ layer is the propagating gauge/matter content actually seen at low energy.

8.5 Backbone constraints carried

Backbone row	How this module carries it
C-FF4 (defining row)	$Q = T_3 + Y$ is *recovered*, row by row, on the §8.6 ledger — the convention's home; Sections 9's table is this one inherited. No $Q = T_3 + Y/2$ variant appears anywhere (§2.4). Certificate: Q02.
C-FF3 (defining row)	This module publishes the five-bundle left-Weyl matter table that Sections 7 and 9 act on — the shared $\mathcal{E}_{\text{matter}}$ in explicit form. Certificate: Q03.
C-FF10	The five anomaly classes close on that same content, with the live per-generation witness $3 \cdot \frac{1}{6} - \frac{1}{2} = 0$ and the even-doublet Witten check (§8.6); detailed authority Paper I Gate 5 / E′. Certificate: Q02.
C-FF0 / C-FF1	Domain is $\mathfrak{B}_{\text{active}}$; carriers S^2 and $S^1_Y/\mathbb{Z}_2$ at ×, projector/charge rules at ⊕, fields at ⊗ (§8.4).
C-FF5	Weak labels and color labels on the same quark are dual roles on one state, per the Appendix C counting rule — never two states. Certificate: Q03.
C-FF8 (supplier side)	The breaking map, the Weinberg rotation, and the surviving $U(1)_{\text{em}}$ are produced here and handed to Section 9 — QED's domain is this module's codomain. Certificate: Q04.
C-FF11	Sanity targets worked in §8.6: the Fermi limit (Section 11, row 3) and tree-level $m_W, m_Z, \rho = 1$. Certificate: Q06.
C-FF12 / C-FF14	Baryogenesis, full precision fits, and full neutrino mechanism are bounded out (§8.8); seven falsification rows live in §8.7.2. Certificate: Q07.

Consumes: the shared Einstein-frame background (Module 6) and Paper I's frozen Higgs/Wilson-line sector (Gate 8: $v = 246.02$ GeV is an *output* there, imported here as a compatibility constraint, never re-derived). **Supplies:** the matter/charge ledger to Modules 7 and 9; the entire electroweak EFT — domain of π_{QED} — to Module 9; the g and G_F rows of §1.2.1.

8.6 The math, simplified then real

Toy first. Break a $U(1) \times U(1)$ with one charged scalar: one combination of the two gauge fields eats the phase and gets a mass; the orthogonal combination stays massless. The electroweak ladder is the non-Abelian version of that single sentence — three generators broken, three Goldstones eaten by $W^{\pm}, Z$, one orthogonal mixture left massless — with the extra structure that the *matter* is chiral, so the same step that splits the boson masses also writes parity violation into the currents.

8.6.1 Weak-sector claim

The weak-force derivation is stricter than simply recovering an $SU(2)$ factor. It must explain chirality, parity violation, symmetry breaking, the massive W and Z, the massless photon, and the Fermi limit. The unified theory therefore treats the weak interaction as the chiral broken descendant of the same 13D metric.

8.6.2 Exact electroweak embedding

The weak-sector projection fixes an exact electroweak embedding inside the active compact geometry. In the 13D routing, weak representations are supplied by S^2 spin-c monopole sectors, while hypercharge labels are supplied by the parent S^1 charge lattice and global $\mathbb{Z}_6$ quotient. The legacy K_6 Cartan formula may still be used as factor-level charge bookkeeping, but it is not the active origin of the weak or hypercharge factors.

The active weak generator is the S^2 spin-c $SU(2)_L$ generator T_3, and hypercharge is quantized as

$$Y \in \tfrac{1}{6}\mathbb{Z}.$$

Electric charge is then

$$Q = T_3 + Y.$$

This is the electroweak embedding used across the unified volume: $SU(2)_L$ from S^2, $U(1)_Y$ from S^1, and the low-energy photon from the unbroken combination $Q = T_3 + Y$.

8.6.3 Branching rules and low-energy chiral matter

The weak-sector projection branches the relevant low irreps under

$$SU(2)_L \times U(1)_Y.$$

The minimal zero-mode assignment used in the unified program is the selected five-bundle SM table, written here in left-handed Weyl notation:

- Q_L: color triplet descendant, S^2 doublet, S^1 charge $Y = 1/6$.
- u_L^c: anti-color triplet descendant, S^2 singlet, S^1 charge $Y = -2/3$.
- d_L^c: anti-color triplet descendant, S^2 singlet, S^1 charge $Y = 1/3$.
- L_L: color singlet, S^2 doublet, S^1 charge $Y = -1/2$.
- e_L^c: color singlet, S^2 singlet, S^1 charge $Y = 1$.

Together with the K_6 index-3 family rule and the active $S^1/\mathbb{Z}_2$ no-mirror lift, this yields the selected three-generation chiral Standard Model spectrum. Parity violation is structural because only the selected left-handed doublet sectors couple to the weak $SU(2)$.

8.6.4 Electroweak anomaly audit

The electroweak anomaly audit closes explicitly. In left-handed Weyl language, the per-generation spectrum is

$$(3,2)_{1/6} \oplus (\bar{3},1)_{-2/3} \oplus (\bar{3},1)_{1/3} \oplus (1,2)_{-1/2} \oplus (1,1)_{1}.$$

For this spectrum, the reduced 4D anomaly coefficients vanish:

$$A_{SU(3)^3} = 0, \qquad A_{SU(3)^2 U(1)_Y} = 0, \qquad A_{SU(2)^2 U(1)_Y} = 0,$$

$$A_{U(1)_Y^3} = 0, \qquad A_{\mathrm{grav}^2 U(1)_Y} = 0.$$

For $SU(2)^2 U(1)$, one generation gives

$$3 \cdot Y(Q_L) + 1 \cdot Y(L_L) = 3 \cdot \tfrac{1}{6} - \tfrac{1}{2} = 0.$$

This is the minimal visible witness that the hypercharge normalization in the table is not arbitrary. This is the one fast-check a reviewer can run with no tools (§0.4 scoreboard): the per-generation sum is identically zero, so quantum consistency is confirmed by a single subtraction.

The global Witten anomaly also vanishes because the total number of left-handed $SU(2)$ doublets is even.

8.6.5 Higgs doublet representation

The electroweak breaking interface requires a Higgs/Wilson-line doublet with

$$H \sim (\mathbf{1}, \mathbf{2})_{1/2}$$

under $SU(3)_c \times SU(2)_L \times U(1)_Y$, using the convention $Q = T_3 + Y$.

8.6.6 Goldstone counting

The breaking

$$SU(2)_L \times U(1)_Y \to U(1)_{\text{em}}$$

removes three generators. Therefore three Goldstone modes are eaten by $W^{\pm}$ and Z, leaving one physical scalar mode in the minimal doublet interface.

8.6.7 Tree-level mass sanity checks

At tree level the interface must reproduce

$$m_W = \tfrac{1}{2} g v, \qquad m_Z = \tfrac{1}{2}\sqrt{g^2 + g'^2}\, v, \qquad \cos\theta_W = \frac{g}{\sqrt{g^2 + g'^2}},$$

and therefore

$$\rho = \frac{m_W^2}{m_Z^2 \cos^2\theta_W} = 1.$$

This paper does not claim a full electroweak precision fit, but failure of these tree-level relations falsifies the weak interface.

8.6.8 Higgs-like mode from geometry

The weak-sector projection derives a light scalar doublet from the compactified geometric sector rather than postulating a separate fundamental Higgs field. That scalar develops a vacuum expectation value and drives the symmetry breaking

$$SU(2)_L \times U(1)_Y \to U(1)_{\text{em}}.$$

The orthogonal combinations become the massive $W^{\pm}$ and Z, while the photon remains massless.

The main GUT manuscript records the Wilson-line/Higgs protection certificate in Paper I App. H (EXTERNAL). This paper imports that status only as a compatibility constraint and does not claim a compactification-independent Higgs theorem.

8.6.9 Weak currents and the Fermi limit

The reduced weak sector furnishes the standard chiral charged and neutral currents. After integrating out the heavy W sector, the low-energy theory reduces to the Fermi interaction, which is the correct short-range limit of the same broken electroweak geometry.

Purpose. Make the low-energy reduction explicit, so a reviewer can read off the four-fermion Fermi interaction from the same electroweak embedding used in the rest of Section 8.

Concretely, integrating out the heavy charged-current exchange yields

$$\mathcal{L}_{\text{Fermi}} \;=\; -\frac{G_F}{\sqrt{2}}\, J^{\mu}_{\text{cc}}\, J_{\mu}^{\text{cc}},$$

with J^{μ}_{cc} the left-handed charged-current built from the same left doublets routed by the S^2 spin-cover. The coupling G_F is the standard Fermi constant set by the W mass and the weak gauge coupling.

This means. At energies far below M_W, the broken electroweak geometry of Section 8 reproduces exactly the four-fermion charged-current interaction of the historical Fermi theory, without introducing any new low-energy operator beyond the declared $V{-}A$ chiral structure.

8.6.10 Link to neutrinos and electromagnetism

The weak-sector projection also serves as the bridge to the neutrino and electromagnetic sectors. Neutrinos live in the same chiral lepton-doublet structure, and the photon is the unbroken Abelian combination that survives electroweak breaking. The downstream QED interface is developed in Section 9.

8.6.11 Formula-by-formula ladder

Each step of the weak-sector construction is anchored on an explicit formula and a short purpose/this-means wrapper, so a reviewer can audit the chain end to end. The seven rungs run from the geometric carrier to the low-energy limit: the weak carrier S^2 (rung 1), its spin cover $SU(2) \to SO(3)$ (rung 2), the hypercharge circle $S^1_Y/\mathbb{Z}_2$ (rung 3), the electric-charge formula $Q = T_3 + Y$ (rung 4), electroweak breaking (rung 5), the photon/Z mixing (rung 6), and the Fermi four-fermion limit (rung 7). Each rung below carries its formula and a plain-language gloss in the same place.

8.6.11.1 Weak carrier

$$S^2 = \{(x,y,z) : x^2 + y^2 + z^2 = R^2\}.$$

Purpose. Fix the compact two-sphere as the weak-sector carrier inside the active 13D routing.

This means. The weak factor is not postulated as an abstract internal $SU(2)$; it is sourced from the geometry of S^2.

8.6.11.2 Spin cover

$$\text{Isom}(S^2) \simeq SO(3), \qquad SU(2) \to SO(3).$$

Purpose. Bridge from the isometry group of S^2 to the weak doublets via the standard spin cover.

This means. Weak doublets are realised on the $SU(2)$ spin cover of $SO(3)$; this is what makes the two-sphere a true weak-isospin carrier rather than just a vector $SO(3)$.

8.6.11.3 Hypercharge

$$S^1_Y/\mathbb{Z}_2, \qquad \theta \sim -\theta.$$

Purpose. Fix the hypercharge direction as a compact circle with the parity identification used in the no-mirror construction.

This means. Hypercharge is geometric, with the $\mathbb{Z}_2$ quotient supplying the boundary/parity data that prevents a mirror partner for each chiral mode.

8.6.11.4 Electric charge

$$Q = T_3 + Y.$$

Purpose. State the electroweak embedding formula that combines weak isospin with hypercharge.

This means. The observed electric charge of every multiplet is the unbroken combination $T_3 + Y$; this is the single formula that the Representation and Chirality Ledger in Section 8.12 must satisfy on every row.

8.6.11.5 Electroweak breaking

$$SU(2)_L \times U(1)_Y \to U(1)_{\text{em}}.$$

Purpose. Fix the breaking pattern so the photon remains massless while $W^\pm$ and Z become massive.

This means. The low-energy QED sector is the unbroken combination, not a separately postulated $U(1)$; the same geometry that supplies $SU(2)_L \times U(1)_Y$ must leave exactly one massless Abelian combination.

8.6.11.6 Gauge-boson mixing

$$A_\mu = \sin\theta_W\, W^3_\mu + \cos\theta_W\, B_\mu, \qquad Z_\mu = \cos\theta_W\, W^3_\mu - \sin\theta_W\, B_\mu.$$

Purpose. Identify the photon and Z as orthogonal mixtures of the W^3 and B gauge fields with weak mixing angle θ_W.

This means. Electromagnetism is the unbroken mixture; Z is the orthogonal mixture and acquires mass from the same electroweak-breaking step.

8.6.11.7 Fermi limit

$$\mathcal{L}_{\text{Fermi}} = -\frac{G_F}{\sqrt{2}}\, J^\mu_{\text{cc}}\, J^{\text{cc}}_\mu.$$

Purpose. State the low-energy four-fermion interaction obtained after integrating out the W exchange.

This means. At energies far below M_W, weak charged-current exchange reduces to the historical Fermi four-fermion interaction; the same chiral left doublets that route the high-energy $SU(2)_L$ also build J^μ_{cc}.

8.6.12 Representation and chirality ledger

The weak-sector matter content is fixed by the following ledger.

Left-handed Weyl field	$SU(3)_c$	$SU(2)_L$	Y	Electric charges
$Q_L = (u_L, d_L)$	$\mathbf{3}$	doublet	$+1/6$	$+2/3, -1/3$
u_L^c	$\bar{\mathbf{3}}$	singlet	$-2/3$	$-2/3$ as conjugate
d_L^c	$\bar{\mathbf{3}}$	singlet	$+1/3$	$+1/3$ as conjugate
$L_L = (\nu_L, e_L)$	$\mathbf{1}$	doublet	$-1/2$	$0, -1$
e_L^c	$\mathbf{1}$	singlet	$+1$	$+1$ as conjugate

The table is in left-handed Weyl notation. Physical right-handed particles are represented by conjugate left-handed fields. This is the convention used for anomaly cancellation.

The hypercharge column is filled in by the active $S^1_Y/\mathbb{Z}_2$ assignment used in Section 8.4, so that $Q = T_3 + Y$ reproduces the observed electric charges on every row.

Required sentence. The weak claim fails if the left/right projector data do not produce left doublets and right singlets with the correct $Q = T_3 + Y$.

8.6.13 Anomaly and no-mirror tests

The electroweak anomaly closure is not optional; it is the binding test that the same surviving matter content used for $Q = T_3 + Y$ in Sections 8.4 and 8.12 also cancels every gauge and global anomaly.

Test	Failure
$SU(2)^2U(1)$ anomaly	electroweak quantum theory inconsistent
$U(1)^3$ anomaly	hypercharge ledger wrong
mixed gravitational-$U(1)$ anomaly	charge assignment wrong
Witten $SU(2)$ anomaly	odd number of left doublets
mirror survival	chiral projection failed

Local pointer. The anomaly ledger is not optional. The weak branch is only acceptable if the same surviving matter content used for $Q = T_3 + Y$ also cancels all electroweak anomalies. The explicit per-generation closure for the spectrum in Section 8.5 is the live witness for this requirement.

8.6.14 Weak-sector result

The weak force is therefore represented as:

- exact chiral $SU(2)_L \times U(1)_Y$ embedding,
- structural parity violation from zero-mode selection,
- geometric Higgs mode and electroweak breaking,
- massive W and Z with massless photon,
- anomaly-free chiral matter content,

- correct low-energy Fermi limit.

8.6.15 The sanity targets and the coupling rows, worked

The C-FF11 targets land on rung 7 and §8.6.9: integrating out the heavy charged current gives $\mathcal{L}_{\text{Fermi}} = -\frac{G_F}{\sqrt{2}} J^{\mu}_{\text{cc}} J^{\text{cc}}_{\mu}$ with $G_F/\sqrt{2} = g^2/8M_W^2$ (Section 11, row 3), and the tree relations $m_W = \frac{1}{2}gv$, $m_Z = \frac{1}{2}\sqrt{g^2 + g'^2}\,v$, $\rho = 1$ are binding falsifiers even though no precision fit is claimed. The §1.2.1 rows: g is the S^2 isometry coupling under the shared normalization, with $\alpha_2(M_Z)$ as Paper I's comparison target; and G_F is **computed** — g from the routing, M_W from v, and v from Paper I's Gate-8 Wilson-line determinant, an output, not an anchor. The weak scale of this paper contains no dial.

8.7 Closure

```
Claim:      the electroweak sector of the active branch is the chiral broken
            descendant of S^2 (spin cover) x S^1_Y/Z2: Q = T3 + Y recovered on
            every multiplet, five anomaly classes closed, EWSB -> U(1)_em with
            massive W/Z, and the Fermi limit as the low-energy reduction.
Mechanism:  spin-cover doublets + hypercharge lattice + no-mirror boundary lift
            + geometric Higgs doublet (1,2)_{1/2} breaking three generators
            (Rungs 1-7; Goldstone count 3).
Why it
matters:    the supplier module: its matter table feeds Sections 7 and 9, its
            breaking is Section 9's domain, and it hosts the routing's live
            arithmetic witness (3*(1/6) - 1/2 = 0).
Paper I
authority:  App. C3 (S^2), C4 (S^1_Y/Z2), D (SM recovery), E/E' (chirality +
            anomalies), H/C9 (Higgs protection; v as Gate-8 output), K
            (neutrinos), R0 (reproducibility).
Certificate: Q02 (charges + anomalies), Q03 (shared bundle homes),
            Q04 (supplier side of composite QED), Q06 (Fermi + tree masses),
            Q07 (scope) — Appendix Q.
Status:     Interface claim (EW recovery + breaking); Main-manuscript
            certificate (chirality, anomalies, Higgs protection, v);
            Outside scope (baryogenesis, precision fits, full nu mechanism).
Falsifier:  wrong Y or Q on any row; mirror survival; any anomaly class
            nonzero; no U(1)_em remnant or wrong mixing; tree relations
            (m_W, m_Z, rho=1) fail; no Fermi reduction (§8.7.2, §13).
```

8.7.1 Minimum reviewer test suite

A minimal acceptance suite for the weak-sector volume is the following.

Test	Required output
S^2 carrier	weak carrier defined
spin-cover bridge	$SU(2) \to SO(3)$ explained
hypercharge carrier	$S^1_Y/\mathbb{Z}_2$ included
charge table	$Q = T_3 + Y$ works for all multiplets
chirality	left doublets / right singlets
no mirrors	mirror modes absent
anomaly	all electroweak anomalies cancel
breaking	$SU(2)_L \times U(1)_Y \to U(1)_{\text{em}}$
Fermi limit	low-energy weak interaction recovered

Sections 8.11–8.13 supply each of these outputs in order.

8.7.2 Falsification / downgrade map

Each weak-sector claim has a declared falsification path and an explicit downgrade if it fails.

Claim	Falsification path	Downgrade
S^2 sources weak sector	no valid spin-cover / weak doublet routing	weak recovery fails
hypercharge consistency	wrong Y or Q for any multiplet	charge interface fails
chirality	mirror modes survive	chirality interface fails
anomaly cancellation	any electroweak anomaly nonzero	anomaly interface fails
EWSB	no $U(1)_{\text{em}}$ remnant or wrong mixing	weak/QED interface fails
Higgs/Wilson-line mechanism	breaking not tied to declared Higgs sector	Higgs compatibility constraint downgrades
Fermi limit	no low-energy charged-current reduction	weak low-energy interface fails

This map is read as a binding requirements list: the entries are conditions under which the interface claim would have to be downgraded, not present failures.

8.7.3 Authority pointers (Paper I)

The downstream evidence for each weak-sector claim lives in the following main-manuscript appendices.

Claim	Authority
S^2	C3 / A1
hypercharge carrier	C4 / D
matter reps	C7 / D
chirality/no mirrors	E
anomalies	E
Higgs breaking/protection	H / C9
neutrino sector	K
reproducibility	R0

These pointers are the audit handles a reviewer can follow from the Section 8 claims back to the canonical authority records in the main GUT manuscript.

8.8 Claim boundary

The boundary below is binding: any statement in this module that appears to overshoot it should be read against this ledger (C-FF12).

8.8.1 Claim boundary ledger

The weak-sector volume operates inside a declared claim boundary. Reviewers should read every subsequent subsection against the following ledger.

Claimed.

- S^2 supplies the weak-isospin carrier structure.
- $S^1_Y/\mathbb{Z}_2$ supplies hypercharge and boundary projection.
- $Q = T_3 + Y$ is recovered for all Standard Model multiplets.
- Left doublets and right singlets are routed through matter/projector data.
- Electroweak breaking leaves an unbroken $U(1)_{\text{em}}$.
- $W^\pm$, Z, and γ are distinguished as low-energy descendants of the same broken electroweak geometry.
- The Fermi limit is recovered as the low-energy charged-current four-fermion interaction.

Not claimed.

- Baryogenesis.
- A complete electroweak precision global fit — **Outside this paper's scope** (OPEN), unless explicitly added.
- The full neutrino mass mechanism beyond the Paper I App. K scope (EXTERNAL).
- Unrestricted Higgs naturalness beyond the declared correction classes.

This boundary is binding: any statement in Sections 8.2–8.15 that appears to overshoot it should be read against this ledger.

---

9. QED / Electrostatic Interface — the Calibration Module

This section is the paper's worked deep instance: the one force module developed at full depth so that the shape of Sections 6–8 can be read at speed. It is the right calibration case for two reasons. First, π_{QED} is the **composite** projection — the only one whose domain is another EFT rather than the active branch — so it exercises the backbone row (C-FF8) that most distinguishes an honest routing from four independent sector notes. Second, its endpoint is the most checkable in physics: Coulomb's law and a massless photon.

9.1 How to read this module (ladder delta)

Two concepts are new here; everything else is assumed from §§3–5 and Paper I.

Composite projection. *Plain:* the photon is not cut directly from the 13D geometry. The weak module first produces the electroweak EFT; QED is then a projection *of that* — the unbroken remnant after breaking. Treating QED as a direct parent projection is not a stylistic alternative; it is the specific error C-FF8 exists to forbid, because it would let the charge convention float free of the weak sector. *Precise:* $\pi_{\mathrm{QED}} : \mathbf{EFT}_{SU(2)_L \times U(1)_Y} \to \mathbf{EFT}_{U(1)_{\mathrm{em}}}$, and the object this module exhibits in formula form is the composite $\pi_{\mathrm{QED}} \circ \pi_{\mathrm{weak}} : \mathfrak{B}_{\mathrm{active}} \to \mathbf{EFT}_{U(1)_{\mathrm{em}}}$.

Static limit as a constraint sector. *Plain:* electrostatics is not a separate classical law sitting beside QED; it is QED with the radiative sector switched off by constraint. *Precise:* the Coulomb interaction is the constrained, non-radiative sector of photon exchange (§9.6.4), so the familiar potential description remains genuinely quantum.

9.2 What QED demands of the routing (plain English)

The low-energy facts any routing must reproduce, stated without formalism: electric charges come in one rigid pattern (the neutron is neutral because quark charges add to zero in exactly the observed combination); the photon is massless and electromagnetism is long-range; charge is locally conserved; the field carries no self-interaction (light does not scatter off light at tree level); and at rest, the force between charges falls as $1/r^2$.

One concrete what-if. If the routing were wrong in the most tempting way — the photon taken as an independent parent projection with its own $U(1)$ — nothing would tie the photon's charge assignments to the weak sector's. The same field could carry $Q = T_3 + Y$ in weak interactions and a freely renormalized charge in electromagnetic ones; atom neutrality would become a coincidence to be tuned rather than an identity to be inherited. The composite structure makes that failure unwritable: QED's charges are the weak sector's charges, by construction of the domain.

9.3 The projection π_{QED}, specified

$$\pi_{\mathrm{QED}} : \mathbf{EFT}_{SU(2)_L \times U(1)_Y} \to \mathbf{EFT}_{U(1)_{\mathrm{em}}}.$$

QED is downstream of the electroweak sector: its domain is the electroweak EFT defined by π_{weak} (Sections 5 and 8), not the active backbone $\mathfrak{B}_{\mathrm{active}}$ directly.

Field	Specification
Projection	$\pi_{\mathrm{QED}} : \mathbf{EFT}_{SU(2)_L \times U(1)_Y} \to \mathbf{EFT}_{U(1)_{\mathrm{em}}}$
Domain	electroweak EFT from π_{weak}
Kept data	unbroken $U(1)_{\mathrm{em}}$, photon, charged matter
Discarded data	massive W/Z at low energy
Projection operation	electroweak breaking + neutral gauge-boson rotation + static limit
Normalization	same $Q = T_3 + Y$ convention as weak section
EFT target	QED action, Maxwell equations, Coulomb law
Sanity check	charge table, Ward identity, photon masslessness, Coulomb law
Falsifier	wrong charges, photon mass term, no Ward identity, no Coulomb limit
Paper I authority	Paper I App. D (SM recovery), Paper I App. H (Higgs/breaking), Paper I App. O (charged-shell audit, diagnostic only) — all EXTERNAL

This instantiates the template of Section 2.5 for the QED sector.

9.4 What survives the projection

Layer	QED/electrostatic content
×	$\mathcal{M}_4$, electroweak compact carriers S^2, S^1_Y
$\oplus$	$Q = T_3 + Y$, $U(1)_{\mathrm{em}}$ selection, gauge-fixing, static-limit rule, conservation rule
$\otimes$	photon field A_μ, charged matter, current J^μ, field strength $F_{\mu\nu}$, stress-energy

The × column lists the carrier geometry. The $\oplus$ column lists the rules that pick out electromagnetism from the broken electroweak sector. The $\otimes$ column lists the fields the rules act on. Discarded at this interface: the massive W/Z sector (integrated out at low energy), heavy KK towers, and modulus modes — discarded means *accounted upstream*, not forgotten; their counting homes are the weak module and the gravity correction ledger (C-FF5).

9.5 Backbone constraints carried

What this module discharges and what it consumes, keyed to the master table of §5.6:

Backbone row	How this module carries it
C-FF8 (defining row)	The domain of π_{QED} *is* the electroweak EFT; the photon is the post-breaking remnant, never an independent parent projection. The composite $\pi_{\text{QED}} \circ \pi_{\text{weak}}$ is the only path from $\mathfrak{B}_{\text{active}}$ to QED in this paper. Certificate: Q04.
C-FF0	Descent from the one parent geometry holds by composition through Section 8.
C-FF1	Every retained object is layer-assigned in §9.4; nothing electromagnetic lives in the × layer.
C-FF3 / C-FF4	The charge table of §9.6.6 is *inherited*, not restated: the same $\mathcal{E}_{\text{matter}}$ assignments and the same $Q = T_3 + Y$ as Section 8, with no convention switch (the §2.4 ledger forbids $Q = T_3 + Y/2$). Certificate: Q02.
C-FF5	The photon is identified with the unbroken combination, not added as a fresh primitive; no mode here is claimed by another sector in an incompatible role (Appendix C, photon row). Certificate: Q03.
C-FF11	The module's sanity targets are worked in §9.6: Coulomb law and the four-row masslessness ledger. Certificate: Q06.
C-FF14	Falsification rows in §9.7; the sharpest is the photon-mass falsifier $m_\gamma^2 A_\mu A^\mu = 0$.

Consumes from Module 8: the breaking pattern, the Weinberg rotation, the Higgs/Wilson-line data behind it, and v. **Supplies downstream:** nothing — QED is the terminal interface, which is why it calibrates the chain.

9.6 The math, simplified then real

Toy first. Take one charged field and one Abelian connection on flat 4D space, and impose gauge invariance alone. Two consequences are immediate: a mass term $m^2 A_\mu A^\mu$ is forbidden, and the coupled current must obey $\partial_\mu J^\mu = 0$. Everything in the ladder below is that two-line statement, descended from the 13D geometry through electroweak breaking instead of postulated.

9.6.1 The descent in prose

After the weak sector breaks to $U(1)_{\text{em}}$, the unbroken massless Abelian combination is identified with the photon. Charge remains geometric because the electric charge operator is still $Q = T_3 + Y$, with Y already fixed by the common electroweak embedding. After integrating out the heavy electroweak, KK, and modulus sectors, the low-energy descendant takes the QED form

$$S_{\text{QED}} = \int d^4x\, \sqrt{-g_4} \left[-\frac{1}{4} F_{\mu\nu} F^{\mu\nu} + \bar{\psi}(i\gamma^\mu D_\mu - m)\psi \right].$$

In the unified program, QED is therefore not fundamental. It is the infrared truncation of the parent 13D electroweak-geometric theory. The electrostatics claim then organizes two things at once: the static limit of the electromagnetic zero mode, and the quantum limit as QED.

9.6.2 Formula-by-formula ladder

> **This ladder ends on Poisson + Coulomb — the static Coulomb/Gauss limit (§9.6.3): the A3 reduces-to-known-physics landing point.** Rungs 8–9 are the cleanest reduces-to-textbook result in the paper.

Each block has a *purpose* statement (what the formula does in the architecture) and a *this means* gloss (plain language).

Rung 1 — Charge operator. $$ Q = T_3 + Y. $$ *Purpose:* fix the electric-charge generator inside the electroweak algebra. *This means:* electric charge is assembled from weak isospin and hypercharge.

Rung 2 — Electroweak remnant. $$ SU(2)_{L} \times U(1)_Y \to U(1). $$ *Purpose:* identify electromagnetism as the surviving unbroken Abelian factor. *This means:* electromagnetism is the unbroken Abelian remnant of the electroweak breaking.

Rung 3 — Photon mixing. $$ A_\mu = \sin\theta_W\, W^3_\mu + \cos\theta_W\, B_\mu. $$ *Purpose:* identify the massless photon as the specific neutral mixture orthogonal to the Z. *This means:* the photon is the massless mixture of the neutral weak field and the hypercharge field.

Rung 4 — QED action. $$ S_{\mathrm{QED}} = \int d^4x\,\sqrt{-g_4}\left[-\frac14 F_{\mu\nu}F^{\mu\nu} + \bar\psi(i\gamma^\mu D_\mu - m)\psi\right]. $$ *Purpose:* fix the low-energy descendant action that the rest of the section must reproduce. *This means:* charged matter couples to the electromagnetic field in the standard 4D QED form.

Rung 5 — Field strength. $$ F_{\mu\nu} = \partial_\mu A_\nu - \partial_\nu A_\mu. $$ *Purpose:* certify the Abelian (commutator-free) structure of the field strength. *This means:* QED is Abelian — there is no self-interaction commutator term.

Rung 6 — Ward identity / charge conservation. $$ \partial_\mu J^\mu = 0. $$ *Purpose:* enforce gauge consistency and local conservation at the level of the current. *This means:* electric charge is locally conserved.

Rung 7 — Maxwell equations. $$ \nabla_\mu F^{\mu\nu} = J^\nu, \qquad \nabla_{[\lambda} F_{\mu\nu]} = 0. $$ *Purpose:* recover the standard sourced and Bianchi equations from the QED action. *This means:* the electromagnetic field obeys the standard source and Bianchi equations.

Rung 8 — Static limit. $$ \nabla^2 \phi = -\rho/\epsilon_0. $$ *Purpose:* reduce the photon zero mode to the scalar potential in the adiabatic, purely electric regime. *This means:* the electric potential is determined by the charge density via Poisson's equation. (Convention: $\mathbf{E} = -\nabla\phi$ and $\nabla\cdot\mathbf{E} = +\rho/\epsilon_0$, so $\nabla^2\phi = -\rho/\epsilon_0$ — the minus sign sits on ϕ, not on Gauss's law.)

Rung 9 — Coulomb law. $$ E(r) = \frac{Q}{4\pi\epsilon_0 r^2}. $$ *Purpose:* close the descent by exhibiting the point-charge solution of the static limit. *This means:* the static Abelian descendant reproduces the ordinary Coulomb field of a point charge.

9.6.3 The sanity target, worked

The C-FF11 target for this module is the pair {Coulomb law, photon masslessness}, and the ladder lands on both. From the covariant Maxwell equation $\nabla_\mu F^{\mu\nu} = J^\nu$, the static electric sector gives

$$\nabla \cdot \mathbf{E} = \rho/\epsilon_0, \qquad \mathbf{E} = -\nabla\phi, \qquad \nabla^2\phi = -\rho/\epsilon_0,$$

and for a point charge $E(r) = Q/(4\pi\epsilon_0 r^2)$ — the electrostatic limit of the same photon field that appears in the QED action (Section 11, row 4). This is the QED sector's reduces-to-known-physics landing point (§0.4 scoreboard): the static limit recovers the textbook Coulomb/Gauss law $E = Q/4\pi\epsilon_0 r^2$ — the same field the textbook gives, recovered, not replaced. (Only the point-charge field and Poisson are worked here; the field-energy/multipole equivalence lives in Paper I App. O, **EXT**. The convention is $\mathbf{E} = -\nabla\phi$ with $\nabla\cdot\mathbf{E} = +\rho/\epsilon_0$, so $\nabla^2\phi = -\rho/\epsilon_0$ — the sign sits on ϕ, not on the field.)

9.6.4 The static limit is a quantum constraint sector

The quantum interpretation is explicit: the Coulomb interaction is not a separate classical law sitting beside QED. It is the constrained, non-radiative sector of photon exchange. The static limit remains genuinely quantum even when it reduces to the familiar potential description.

9.6.5 Photon masslessness ledger

The masslessness of the photon is tracked as an explicit pass/fail ledger, not as an assumption.

Requirement	Pass condition
$U(1)_{\text{em}}$ unbroken	photon has no mass term
charge generator preserved	$Q = T_3 + Y$ commutes with vacuum / Wilson-line sector
no boundary mass	orbifold / boundary terms do not generate a photon mass
Ward identity	charge conservation holds

All four rows must pass for the QED interface to close. The falsifier is the survival of any photon mass term: $m_\gamma^2 A_\mu A^\mu = 0$ is required; if electroweak breaking, boundary terms, or compactification corrections generate a nonzero photon mass term, the QED interface fails.

9.6.6 Charge normalization table

This volume uses $Q = T_3 + Y$ with the Y values shown below. The four left-handed fields fix the normalization of the entire matter sector.

Field	T_3	Y	$Q = T_3 + Y$
ν_L	$+1/2$	$-1/2$	0
e_L	$-1/2$	$-1/2$	-1
u_L	$+1/2$	$+1/6$	$+2/3$
d_L	$-1/2$	$+1/6$	$-1/3$

The remaining matter charges follow from this normalization through the standard electroweak doublet/singlet structure.

9.6.7 The coupling row of §1.2.1, worked

The derived coupling for this module: $e = gg'/\sqrt{g^2 + g'^2}$, read off Rung 3 — the photon's coupling is the normalization of the massless mixture, computed from the weak and hypercharge couplings routed in Sections 8 and 5. No new dial enters; this is the table row that makes the five-couplings headline falsifiable at the QED interface (C-FF4/C-FF8, Section 13).

9.7 Closure

```
Claim:      QED with its electrostatic limit is the unbroken Abelian descendant of
            the electroweak interface — the composite π_QED ∘ π_weak — with charges
            Q = T3 + Y inherited, photon massless by ledger, Coulomb/Gauss recovered.
Mechanism:  electroweak breaking + neutral-boson rotation + static limit (Rungs 1–9).
Why it
matters:    the terminal interface; calibrates the whole routing — if the composite
            structure failed here, "one geometry, four forces" would already be false.
Paper I
authority:  App. D (SM recovery), H (breaking), C3/C4/C7/C8 (dossiers), O (shell
            audit, diagnostic), R0 (reproducibility), R2 (claim vocabulary).
Certificate: Q02 (charges), Q03 (counting homes), Q04 (composite factorization),
            Q06 (sanity targets) — Appendix Q.
Status:     Interface claim (the routing); Main-manuscript certificate (every
            invoked gate); Diagnostic (shell audit, §9.8.3); Outside scope (§9.8).
Falsifier:  photon mass term; wrong charge table; no Ward identity; Coulomb limit
            fails; or QED written as an independent parent branch (§9.7.2, §13).
```

9.7.1 Minimum reviewer test suite

Test	Required output
charge operator	$Q = T_3 + Y$
photon origin	$SU(2)_L \times U(1)_Y \to U(1)_{\text{em}}$
photon masslessness	no photon mass term
QED action	standard kinetic + charged matter coupling
Ward identity	$\partial_\mu J^\mu = 0$
Maxwell equations	recovered
static limit	Poisson equation
point charge	Coulomb law
scope	shell audit illustrative only

Each row corresponds to a specific formula or boundary statement in §9.6 and §9.8.

9.7.2 Falsification / downgrade map

Claim	Falsification path	Downgrade
EM is electroweak remnant	no unbroken $U(1)_{\mathrm{em}}$	QED interface fails
charge rule	wrong charge table	charge interface fails
photon massless	photon mass term appears	QED fails
QED action recovered	wrong kinetic/coupling structure	QED descendant fails
Ward identity	charge not conserved	gauge consistency fails
electrostatic limit	Poisson/Coulomb not recovered	electrostatic interface fails
shell audit	mismatch with Maxwell accounting	shell-audit diagnostic fails
shell audit overused	used as GUT proof (the shell example is illustrative only, **Outside this paper's scope** as a gate; status OPEN)	claim-boundary failure

9.7.3 Authority pointers (Paper I)

Claim	Authority in Paper I (Fable_GUT) — all EXTERNAL
charge operator	App. D / dossiers C3, C4
EW breaking	App. H / D
photon field	App. D / H / dossier C8
charged matter	dossier C7
QED equations	dossier C7 / App. O
shell audit	App. O
reproducibility	App. R0 (incl. R0.R certificate set)
claim boundary	App. R2

9.8 Claim boundary and scope annexes

Claimed: $U(1)_{\mathrm{em}}$ is the unbroken electroweak descendant; electric charge is $Q = T_3 + Y$; the photon zero mode is massless within the declared breaking mechanism; the QED action is recovered at low energy; the Ward identity / charge conservation is respected; electrostatics is the static limit (Gauss, Poisson, Coulomb).

Not claimed: a full nonperturbative QED theorem package; condensed-matter electrostatics or triboelectricity as a fundamental-physics prediction; all radiative corrections beyond the declared renormalization scope; the charged-shell example as required GUT-gate closure.

9.8.1 Standard objections, answered at interface level

The Landau pole. This paper does not solve the QED Landau-pole problem. It treats standalone QED as the low-energy Abelian EFT obtained after electroweak breaking. Extrapolating isolated QED beyond its effective domain is outside this paper's claim.

The unified-theory gap. QED by itself does not include gravity. In the unified theory, this gap is addressed at model-definition level because electromagnetism and gravity are different projections of the same 13D metric.

The static mystery. The longstanding difficulty of giving a universal microphysical account of all static charging phenomena is handled, as an optional downstream application only, by the support-field/kernel description collected in §9.8.4; it is not used to close the QED/electrostatic interface claim of this section.

9.8.2 Renormalization scope note

QED radiative corrections are standard EFT corrections after the low-energy QED action is recovered. The force-interface document does not need a full nonperturbative QED completion to recover electrostatics.

9.8.3 Conservation-audit example boundary

The charged-shell audit must say:

> The charged-shell example is illustrative same-output equivalence. It demonstrates that the geometry ledger reproduces standard Einstein-Maxwell accounting for a specified configuration. It is not used to close any required GUT gate.

9.8.4 Optional downstream application: contact charging and triboelectric kernels

The following material is an optional downstream application of the QED/electrostatic interface. It is not required for the four-force routing claim and is not used to support any GUT gate certificate.

Surfaces and materials are represented by localized supports, channel kernels, and overlap-generated effective coefficients; contact charging is then modeled as the approach of two support sectors, overlap-driven charge transfer through a contact kernel, trap occupation and internal redistribution, and geometric freeze-out when the surfaces separate. This provides a condensed-matter / materials-level account of triboelectric charging that sits inside the same geometry-to-effective-physics architecture as the rest of the interface, but it is illustrative engineering content only: it is condensed-matter modeling, not a fundamental-physics claim, and it is explicitly listed above under "Not claimed" as a fundamental-physics prediction.

---

10. Cross-Sector Consistency Conditions

The cross-sector consistency conditions below are the downstream audit of the Four-Force Constraint Backbone in Section 2.6. Section 2.6 states the constraints; this section checks whether the four instantiated force projections satisfy them together.

Each force interface defined in Sections 6–9 must obey shared conventions for counting states, cancelling anomalies, and assigning regulators across the four sectors. The three ledgers below collect those shared conventions in one place so that a sector-by-sector reading cannot accidentally introduce a convention conflict. Any conflict triggered by these ledgers downgrades the affected interface claim via the falsification map in Section 13.

10.1 State-counting ledger

Every physical state used by one force must be assigned to exactly one sector for counting purposes. Double-counting a state across two sectors or omitting it from all sectors falsifies the integration.

Sector	States counted	Shared authority
gravity	metric modes after gauge quotient	gravity interface
strong	gluons + quark color states	matter/gauge bundles
weak	chiral doublets/singlets	matter/gauge bundles
QED	charged matter + photon	matter/gauge bundles
flavor	family states	flavor chamber

Failure statement. If the same state is counted in two sectors, or if a state is omitted from all sectors, the integration fails and the affected interface claim is downgraded.

10.2 Anomaly ledger

The anomaly content of the unified theory must close across the four sectors using a single matter content. The matter content used to recover $Q = T_3 + Y$ in the weak/QED interface must also cancel the anomaly classes listed below.

Anomaly class	Relevant sectors
gauge anomalies	strong / weak / hypercharge
mixed gauge-gravity	all charged matter
boundary/orbifold anomalies	weak / hypercharge / QED
BRST/gauge anomalies	quantum companion

10.3 Regulator/RG ledger

Threshold corrections, running couplings, and charge normalizations must use compatible conventions across sectors. Inconsistent conventions falsify the cross-sector integration even when each sector is internally consistent.

Use	Required consistency
threshold unification	same compact spectrum and regulator convention
QCD beta function	compatible state counting
QED running	compatible charge normalization
electroweak running	compatible $SU(2)/U(1)$ normalization
quantum-gravity companion	not used to close interface claims unless separately routed through the main manuscript

10.4 Cross-Sector Convention Matrix

Convention	Gravity	Strong	Weak	QED
4D metric	$g_{\mu\nu}$	background $g_{\mu\nu}$	background $g_{\mu\nu}$	background $g_{\mu\nu}$
matter notation	stress-energy	Dirac for action / Weyl for anomaly	Weyl for reps/anomaly	Dirac for low-energy action
charge convention	n/a	color only	$Q = T_3 + Y$	same $Q = T_3 + Y$
gauge normalization	n/a	$\mathrm{Tr}(T^aT^b) = \frac{1}{2}\delta^{ab}$	declare $SU(2), U(1)$ normalization	inherits EW
scope exclusions	QG	confinement / strong CP	baryogenesis / full precision	nonperturbative QED / contact charging

10.5 Interface Consistency Proposition

Proposition. Given the active branch $\mathfrak{B}_{\text{active}}$, the Global Conventions Ledger of Section 2.4, the Four-Force Constraint Backbone of Section 2.6, and the four Operational Projection Specifications of Sections 6–9, the four-force interface is well-defined if and only if:

1. each π_i has a declared domain, codomain, kept data, discarded data, normalization, and EFT target;
2. all four projections use the same matter-bundle and charge conventions;
3. no state, generator, or field is counted in two incompatible sectors;
4. anomalies and regulator conventions are shared across sectors;
5. QED factors through electroweak breaking;
6. every force section has at least one low-energy sanity target.

This is an Interface claim — an interface proposition, **CONDITIONAL** on its premises, not a proof of final physical truth.

10.6 Interface Falsifier

The proposition is false if any force projection is undefined, if two projections assign incompatible quantum numbers to the same matter field, if QED is not downstream of electroweak breaking, if the same mode is counted twice in incompatible ways, or if a force section lacks a recognizable low-energy EFT target.

10.7 Regulator / RG compatibility rule

Any beta-function, threshold, or running-coupling statement must use the same state-counting convention and generator normalization declared in the Global Conventions Ledger. If the QCD, electroweak, QED, and threshold sections use incompatible normalizations, the cross-sector interface fails even if each sector looks locally correct.

11. Observable Sanity Targets

A minimum-pass sanity check per force is a recognizable low-energy or 4D limit. The targets below are necessary; the main GUT manuscript provides the corresponding gate certificates. This is the C-FF11 sanity-target object (Q06 reads its five frozen formulas verbatim, and the modules' §N.6 ladders land on its numbered rows). The honest grade for each landing point is the §0.4 scoreboard — not re-printed here; this table adds only the two checks §0.4 does not carry: the strong-sector asymptotic-freedom **sign**-check (row 2) and the EW/QED photon-masslessness ledger (row 5).

Force	Minimal target	Formula / check	Status
Gravity	Newtonian limit	$\nabla^2\Phi = 4\pi G\rho$	grade per §0.4 (worked §6.6.9)
Strong	asymptotic freedom	$\beta(g_s) < 0$ for SM active flavors	Interface claim; consistency check (sign only), confinement scoped out (§7.8)
Weak	Fermi limit	$G_F/\sqrt{2} = g^2/(8M_W^2)$	grade per §0.4 (derived identity, §8.6.15)
QED	Coulomb law	$E = Q/(4\pi\epsilon_0 r^2)$	grade per §0.4 (worked §9.6.3); recovers the static Maxwell limit, does **not** replace QED/Maxwell
EW/QED bridge	photon masslessness	no $m_\gamma^2 A_\mu A^\mu$	Interface claim; massless by ledger (§9.6.5)

12. What This Paper Does Not Claim

The authoritative non-claims list is **§2.2** (stated early, where scope belongs); every item there is **Outside this paper's scope** (status: OPEN) — none is established, refuted, or upgraded here. This section adds only the two boundary items §2.2 does not spell out:

- **The QCD mass gap — Outside this paper's scope** (OPEN); not established or refuted here (the strong sector's confinement/χSB edge is open, §7.8, WL-4).
- **All precision phenomenology beyond the leading observable targets of Section 11** — the per-sector sanity targets there are necessary minimal checks, not full precision fits.

Companion or audit material that may exist for these topics outside the main GUT manuscript is not used to upgrade the main manuscript's gate certificates, and is outside this paper's scope.

13. Falsification and Downgrade Map

Each interface claim of this paper has an explicit falsification path and a downgrade consequence. If any row triggers, the affected interface claim and any downstream consequence in Section 10 are downgraded. Each row is tied to a named constraint from the Four-Force Constraint Backbone of Section 2.6, so a reviewer can locate the first failing link directly; the new Certificate column names the runnable Appendix Q check for each row.

Constraint	Integration claim	Falsification path	Downgrade	Certificate
C-FF0	one geometry	any branch uses different parent geometry	unified interface fails	Q05
C-FF2	projection-based forces	any force lacks a projection map	branch becomes external insertion	Q05
C-FF5	no double-counting	state/field counted twice	affected interface claim downgrades	Q03
C-FF4	shared charge convention	inconsistent Q normalization	weak/QED interface fails	Q02
C-FF9	regulator compatibility	incompatible RG/threshold schemes	cross-sector RG claims downgrade	Q06
C-FF10	anomaly compatibility	any anomaly class unaccounted	weak/QED interface fails	Q02
C-FF8	QED downstream of EW	photon not post-breaking remnant	weak/QED interface fails	Q04
C-FF12, C-FF13	companion claims scoped	quantum-gravity or nonperturbative-QCD companion used to upgrade an interface claim	scope-boundary failure	Q07
C-FF1, C-FF6, C-FF7	layer / gauge-origin / metric-block discipline	gauge field smuggled into $\times$ layer, or metric/mixed/compact blocks conflated	layer-discipline failure	Q03/Q05
C-FF3	shared matter bundle	different forces act on incompatible $\mathcal{E}_{\text{matter}}$	matter-content interface fails	Q03
C-FF11	low-energy sanity target	a force section lacks a recognizable 4D EFT target	sanity-target interface fails	Q06
C-FF14	falsification-path discipline	any projection or integration claim lacks a stated falsifier	hostile review cannot localize failure; integration claim downgrades	Q01

All 15 backbone rows C-FF0–C-FF14 appear above (grouped where one falsifier covers several — e.g. C-FF12/C-FF13, and C-FF1/C-FF6/C-FF7); none is unmapped.

13a. The Strongest Objections, Answered

Plain. A serious reviewer arrives with serious objections, and an interface paper that ducks them earns nothing. So here are the four strongest objections a hostile expert raises against a routing like this one — each stated **at full force first**, in the form its sharpest critic would use, and only then answered, by pointing at the exact place in this paper where the claim → mechanism → certificate → falsifier chain already lives. No objection below is softened before it is answered; where an answer is partial, it says so. This section makes **no new claim and changes no status** — it is a reader aid that routes each objection to material already in the body.

Objection 1 — "Four projections of one geometry is just relabeling. You can always carve four EFTs out of a big enough manifold and call the carving 'unification.'" *Stated at full force* (the objection's claim, status OPEN until answered)*:* writing four maps π_i out of one space proves nothing if each map is free to discard whatever it likes; the "unification" is then bookkeeping, and the four sectors never actually constrain one another. *Answered:* the constraint is precisely that the four maps are **not free** — every π_i must survive the *same* fifteen-row Four-Force Constraint Backbone (C-FF0–C-FF14, §2.6), and the binding ones are *joint*: one matter bundle acting in all four sectors (C-FF3), one charge convention $Q = T_3 + Y$ surviving the weak→QED composition (C-FF4, C-FF8), and a no-double-counting ledger that closes over the *union* of all four state inventories (C-FF5; §10.1). Any single sector's carving is blind to these joint meshes, so passing them is not free. A relabeling cannot pass a joint mesh it was never built to satisfy — the by-hand electroweak anomaly witness (§8.6.4) is the example of a mesh the routing must hit and an arbitrary carving would miss.

Objection 2 — "Zero new dials is a slogan. You imported the dials from Paper I and are pretending the bill was already paid." *Stated at full force:* the economy headline only holds if the upstream inputs were genuinely few and fixed before comparison; if Paper I quietly tuned a dozen knobs, "zero new dials *here*" is an accounting trick that hides the fitting one paper upstream. *Answered:* correct that the economy is **inherited, not re-proved here** — and the paper says exactly that, routed to the top as a declared weakest link (WL-3, §0.2: all gate certificates are external). The "zero new dials" claim this paper *owns* is narrower and checkable: the only measured numbers in the two-paper system are Paper I's **four** declared inputs ($M_{\rm Pl}$; $\alpha_i^{-1}(M_Z)$; y_t; $|V_{us}|$), and the inputs-in/outputs-out ledger (§1.2.1) shows the five force strengths flowing out of that same economy with **no further input added here** (CALIBRATED anchors and comparison targets, labeled). Two of the five — e and G_F — are not imported at all: they are **computed** from the others ($e = gg'/\sqrt{g^2 + g'^2}$; $G_F/\sqrt{2} = g^2/8M_W^2$), and they are *identities of the routing*, not predictions (§1.2.1, §9.6.7, §8.6.15). If Paper I's input economy ever downgrades, the consuming rows here downgrade automatically and one-way (§14.3) — the bill is not hidden; it is itemized and cross-linked.

Objection 3 — "Calling e and G_F 'derived' is overclaiming. Every electroweak textbook writes those relations; you've rediscovered the Standard Model, not derived it from geometry." *Stated at full force:* $e = gg'/\sqrt{g^2 + g'^2}$ and $G_F/\sqrt{2} = g^2/8M_W^2$ are standard electroweak identities. Presenting them as a geometric achievement inflates a textbook result into a discovery. *Answered:* the paper makes the **weaker, honest** claim, and labels it as such: e and G_F are *derived identities of the routing*, **not predictions** (§1.2.1, A10) — the contribution is not the relations themselves but that they fall out of *this* geometry's electroweak descent with no independent dial, so that they become **falsifiable teeth** rather than free parameters. The geometric content is specifically that g' comes from the $S^1_Y/\mathbb{Z}_6$ hypercharge lattice and g from the S^2 spin cover (§8), at fixed normalizations; if either identity failed against measurement at those normalizations, the weak/QED interface would be wrong *as an interface*, independent of Paper I (§13, rows C-FF4/C-FF8). The honest boundary is stated verbatim in §12: this paper claims **no new numerical prediction beyond Paper I's**. Naming a relation "derived" here is licensed only because the derivation and its falsifier both sit in-doc — never as a bare promotion.

Objection 4 — "The hard parts are all scoped out. Confinement, strong CP, quantum gravity, the cosmological constant — you've excluded exactly the questions that are hard, then claimed the easy remainder as a triumph." *Stated at full force:* an interface that recovers $SU(3)_c$ but punts on confinement, recovers a gravity EFT but punts on UV completion, and waves off strong CP and Λ has answered the tractable questions and renamed the intractable ones "out of scope." *Answered:* the scope boundary is real, and the paper's discipline is that a scoped-out sector is **bounded, never used as support** — this is scope-boundary-*as-strength*, recorded as a binding backbone row (C-FF12; §2.2, §12). The honesty is mechanical: nonperturbative QCD, strong CP, full quantum gravity, and cosmology/Λ each appear only with a status label and a pointer to where they live upstream (the strong-CP and confinement edge at §7.8, openly underived; the stabilization and Λ questions cross-referenced to the corpus capsules at §6.8.9, EXTERNAL), and the no-double-counting and regulator ledgers (§10) forbid any of them from feeding an interface claim. Crucially, the falsifier survives the exclusion: the falsification map (§13) keys every interface claim to a backbone row, so a reviewer can localize the first failing link *without* the scoped-out physics. A paper that claimed those sectors would be the one to distrust; one that recovers the gauge *interface* at a stated grade and routes the deep dynamics to their upstream owners, with a falsifier on every retained claim, is showing its evidence chain.

The takeaway is: none of these four objections is answered by assertion; each is answered by a pointer to a backbone row, a runnable Appendix Q certificate, a named upstream authority, or a pre-registered falsifier already in this paper. That is the difference between an interface asking to be believed and one inviting you to check.

---

14. Relation to Paper I: the Interface Ledger

This section is the formal boundary between this paper and Paper I (*Fable_GUT*): the provenance note and the legacy-reference convention (§14.1), the authority map giving Paper I's home for every item this paper invokes by interface (§14.2), the one-way cross-paper downgrade-inheritance rule (§14.3 — if a Paper I gate downgrades, the consuming rows here downgrade automatically, and the converse never holds), and a statement of what this paper adds that Paper I does not contain (§14.4). Every "frozen", "certificate", "anomaly-free", or gate-status claim in this paper resolves through this ledger to Paper I, EXTERNAL; this paper re-closes nothing.

14.1 Provenance

This paper is the standalone force-interface companion to the main GUT manuscript, **Paper I: *Fable_GUT*** — the restructured build organized around the gate–constraint inversion (its §1.3), with the worked two-anchor fixing as its Section 8 and the machine-certificate registry at its R0.R. An earlier draft of Paper I circulated as *final_manuscript* (https://physics.magflowmeters.com/gut/manuscript/final_manuscript.html); all authority letters below resolve in the Fable_GUT build, whose A3.R ledger maps any legacy reference. This paper uses the same active geometry and notation and reproduces none of Paper I's certificates.

14.2 Authority map

Paper I remains the sole authority for the items this paper invokes by interface:

Invoked here	Paper I home (all EXTERNAL)	Used by
Frozen geometry and reconstruction	App. A1 / A2; Paper I §2	§§3–4; every module §N.3 domain row
Three-layer rule	Paper I §2B; App. B2	§4; routing tables §N.4
Term-level necessity dossiers	App. C (C1–C9)	module authority tables §N.7.3
SM recovery; charges	App. D; Gates 2–3	Modules 7–9
Chirality and anomaly closure	App. E / E′; Gates 4–5	Modules 7–8 (Weyl ledgers)
Moduli freeze / witnesses	App. F; Gate 6	Module 6 (Einstein-frame caveat)
Threshold unification	App. G; Gate 7	Module 7; §1.2.1 (g_s, g targets)
Higgs protection; v, m_h	App. H; Gate 8; Paper I §8	Module 8 (breaking; G_F row)
Charged-shell field-energy equivalence (worked)	App. O (O.4 / O.5)	§9.6.3, §9.8.3 (this paper carries the diagnostic-only boundary); §0.4.1 Prompt 1 (EXT)
CKM CP phase δ_{CKM} (raw $-120°$ → Wolfenstein $+60.0°$)	§7.6; App. I / J	§0.4.1 Prompt 4 (EXT; scoped out here, §2.2)
Generation count (index -3); LEP N_ν consistency	§1.3.1; §3.4–§3.5; App. E	§0.4.1 Prompt 5 (EXT)
Flavor chamber and outputs	App. I / J / K; Gate 9	quark masses (Module 7); §10.1 flavor row
Proton safety	App. L; Gate 10	scope of ⊗ operator content
Claim-boundary protocol	Paper I §9; App. R2; Gate 11	§2.2, §12, module §N.8 annexes
Freeze records, hashes, reproducibility	App. R0 (incl. R0.R), R1	every "frozen" claim in this paper

14.3 Cross-paper downgrade inheritance

The dependency edge is one-way and binding: **if a Paper I gate downgrades, the rows here that consume it downgrade automatically**, with no renegotiation in this paper. The inheritance map: Gate 2/3 → C-FF4 rows (Modules 7–9 charge tables); Gate 4/5 → C-FF10 rows (Modules 7–8 ledgers); Gate 6 → Module 6's Einstein-frame claim (drops to scalar-tensor); Gate 7 → the g_s/g comparison rows of §1.2.1; Gate 8 → Module 8's breaking scale and the G_F row; Gate 9 → the quark-mass descendant claims of Module 7. The converse never holds: nothing in this paper upgrades, re-closes, or substitutes for a Paper I certificate (C-FF13).

14.4 What this paper adds

A different object from Paper I: a unified routing map — the four projections, their backbone, the cross-sector ledgers, the coupling-provenance table — showing gravity, strong, weak, and QED/electrostatics as mutually compatible reductions of one 13D geometry, with its own machine-certificate layer (Appendix Q) checking the interface claims that are this paper's contribution.

---

15. Conclusion

Detailed GUT gate certificates remain in the main manuscript; this paper supplies the interface layer. The four-force interface is therefore not merely a set of parallel reductions. It is a constraint-filtered routing statement: one parent geometry, one layer discipline, one matter/charge convention, one no-double-counting rule, one regulator/anomaly discipline, and four force-specific projections. The Four-Force Constraint Backbone of Section 2.6 makes explicit the conditions under which the gravity, strong, weak, and QED/electrostatic sectors can be read as mutually compatible projections of the same active branch. If any backbone constraint fails, the affected interface claim downgrades according to the falsification map of Section 13.

15.1 Backbone coverage

Backbone coverage in both tenses — every C-FF row enforced by a named module and checked by a runnable Appendix Q certificate — is the §5.6 master table (the authoritative home for the coverage claim).

Stated in this build's terms: the Four-Force Constraint Backbone is the paper's one list in two tenses — the conditions that defined the four projections are the conditions the finished interfaces keep passing, with a runnable certificate per row — and the headline that list buys is §1.2.1: five force strengths, one geometry, zero new dials.

There is a way to test whether a geometry is real that has nothing to do with particle physics: **try to build something with it.** A shape that is mere bookkeeping cannot be machined into a working device; a shape that is real can — and §16 records the same coset, built and simulated.

Read on: §16 carries that bridge in full, at its honest engineering grade — the admissibility chamber of a quantum error-corrector, designed and simulated on open tools.

---

16. The Geometry Is Not Abstract — The Same Coset Whose Isometries Are the Forces Builds a Quantum Computer

> **Status of this section, for the record.** An EXTERNAL engineering bridge at **SIMULATED** (engineering) grade. It makes **no new interface claim** and changes **no status** in §§5–13: every projection, coupling row, and falsifier stands exactly as the body records it. Zero promotions. The QC corpus (the Tahoe Labs error-correction design package and its `SIMULATION_CAMPAIGN` / `GATE_CLOSURE` evidence) is EXTERNAL and registered here once; nothing in it is reproduced or modified by this document, and nothing in it moves a single interface row. No geometry is changed (MR-6 / C-FF0 geometry-change test: N/A — this section *reports* a pre-existing engineering artifact and alters no manifold, coset, radius, code distance, or lattice in this paper; the physical:logical qubit ratio at 10,000 logical qubits is therefore untouched). This section is byte-consistent with the engineering-bridge sections of the sibling capsules TOE_FINAL_merged and Fable_Quantum_merged (UQFC-19.5) (EXTERNAL) on every quantitative figure below.

Plain. Every section above asks you to believe one thing: that a particular six-dimensional shape — $K_6 = SU(3)/T^2$, the squashed flag manifold of the active branch — is *physically real*, the actual internal geometry of the world rather than a convenient piece of mathematics. The whole point of this paper is that the **isometries of that shape are the forces**: the strong $SU(3)_c$ is the surviving left action of K_6 (§7.3), and the weak and electromagnetic sectors descend from the companion S^2 and S^1_Y factors of the same frozen geometry (§8–§9). There is a way to test whether a geometry is real that has nothing to do with particle physics: **try to build something with it.** A shape that is merely a bookkeeping device cannot be machined into a working device; a shape that is real can. The same $K_6 = SU(3)/T^2$ coset — written in the QC corpus in its maximal-torus form $SU(3)/U(1)^2$, the identical manifold (Appendix A: $T^2 \cong U(1)^2$) — defines the *admissibility chamber* of a quantum-error-correction architecture that has been designed and simulated on standard open-source tools. This section records that bridge at its honest engineering grade — neither more nor less — because a force-interface theory whose geometry earns its keep in a buildable machine has crossed from formalism into engineering.

Precise. The bridge is structural, not metaphorical — and the receipt for that is that it has been run, not just described: the same admit-one-eigenspace/reject-the-rest operation is implemented as runnable code (§16.2). It makes **no new claim about the four-force interface** and changes **no status in §§5–13**. The routing of this paper admits the physical force content out of the K_6 isometry algebra and rejects the rest — the surviving $SU(3)_c$ generators are *admitted*, the broken/projected-out directions are *rejected* (the layer discipline of $\times/\oplus/\otimes$, §3–§5). The error-correction architecture is built on the identical operation in a different register: a **coset-constrained admittance algebra** that admits only the $+1$ eigenspace of a Heisenberg–Weyl / holonomy stabilizer subgroup and rejects its complement. *Admit one eigenspace of a coset-defined operator, reject the rest* is the same mathematical move in both places. The QC program is a **sibling engineering corpus** (EXTERNAL, registered below); nothing in §16 feeds back into the interface claims, and the no-target-loading discipline (C-FF13: gate authority is invoked by interface only, never imported as new structure) forbids any such feedback.

16.1 The shared object: one chamber, two uses

The same structural idea	In this paper (force routing, §5–§9)	In the QC architecture (sibling corpus, EXTERNAL)
The host geometry	$K_6 = SU(3)/T^2$ flag manifold — the active internal space (§2.1, Appendix A)	The $K_6 = SU(3)/U(1)^2$ chamber of the code's stabilizer layer (the same coset; the design's "chamber")
The admit/reject operation	Layer discipline admits the surviving $SU(3)_c$ isometry action and the descended electroweak generators, rejects the projected-out directions (§5.1, §7.3)	Admit the $+1$ eigenspace of the stabilizer subgroup, reject the complement (the admittance algebra)
What it controls	Which gauge forces are physical — the strong/weak/EM projection maps (§7–§9)	Which logical states are protected — the code's admissible subspace
The selectivity number	The surviving generator count per sector (8 gluons; 3 weak $+1$ hypercharge $\to W^\pm/Z/\gamma$) that each projection must reproduce	The chamber's admitted-fraction / selectivity (the figure that drives the qubit budget)

The *same coset and the same admit-one-eigenspace-reject-the-rest operation* appear in both columns; in the QC case that operation has been written as runnable code and exercised against a decoder (§16.2).

16.2 What has actually been built and simulated (honest engineering grade)

The following are recorded at **engineering grade** — the grade of a *simulated design*, not of a fabricated and measured chip, and explicitly *not* at this paper's interface-certificate grade. Each carries its status label:

- **The chamber concept is validated in simulation (SIMULATED).** The admissibility-chamber projector is implemented and exercised; the four-layer error-correction stack it sits in (GKP + chamber + surface + modular) runs end-to-end, and the simulation suite executes cleanly — `verify_campaign.py` exits 0 with its core checks PASS on live re-runs (SIMULATED; EXTERNAL, the QC corpus).
- **The design is locally simulatable today (SIMULATED).** The toolchain is open-source and installed — Stim + PyMatching (stabilizer simulation and decoding), QuTiP and scqubits (the bosonic / transmon layer), with EM parasitic extraction verified. The architecture is not slideware; it is exercised against a real decoder at ~10^8 shots (SIMULATED).

- **The honest overhead floor is competitive, not heroic (BOUNDED).** The design's physical-to-logical memory overhead at scale sits in a **~9.6–14.5:1** band (anchor 14.5:1, frontier 9.6:1; BOUNDED, SIMULATED), with a **contingent robust four-layer fallback of 3.24:1** (BOUNDED) — competitive with published qLDPC overheads (~24:1). The conservative number is the anchor 14.5:1, not the 9.6:1 frontier; the 3.24:1 fallback is the smaller four-layer figure and is contingent, not a headline. These are competitive numbers for a fault-tolerant *memory*; the formerly-quoted sub-2 ratios are **retired** — they were an algebra-clamp artifact and a single-module figure, withdrawn on the record (see the self-caught discipline below). No sub-2 ratio is claimed anywhere in this document.
- **The decoder advantage is real and survived audit (SIMULATED).** The biased-noise advantage of the deformed code was tested with a genuine Hadamard-deformed XZZX circuit in Stim with a deterministic observable; the design operates **below threshold at its operating point** (SIMULATED), with a **simulated** single-logical biased advantage of ~8.4× (CI 4.2–16.9; SIMULATED) on the decoded logical channel.

16.3 Why the engineering discipline is itself the evidence

The reason to trust the bridge is the same reason to trust the physics in this paper: **the QC program caught and reversed its own optimistic errors, and got more credible for it** — exactly the discipline this paper applies when it leaves its most tempting edge open (strong CP / confinement: verified out-of-scope and *not* used as support, §7.8). The honest-discipline record carries at least **seven self-caught corrections** (SIMULATED / BOUNDED grades as marked); the two with the clearest stakes are shown in full, the other five collapsed:

1. **A raw-syndrome "4.3×" decoder gain was REVERSED** as an undecoded artifact when a real decoded simulation was run; the honest *decoded* figure — a different, properly-decoded quantity — is ~8.4× (CI 4.2–16.9), consistent with the original conservative *decoded* claim (SIMULATED). The two numbers measure different channels, so the move is a correction, not an upgrade.
2. **A $d \geq 8$ classical-signaling claim was DOWNGRADED** (REFUTED-class) to PAM-8-plus-parity once a computation (with the `galois` library) proved that the construction was the kernel of a GF(2) parity check — the coset advantage survives only in the coherent quantum embodiment, and the design says so (recorded).
3. **Five further self-corrections are logged in the QC corpus (EXTERNAL):** a thermal-via benefit lowered 15–25% → ~6.5% (FEM spreading resistance); an SU(N) selectivity escalator lowered when the admitted fraction was shown to scale as $1/(N+1)$ — the inconvenient direction; a success probability corrected from a 91% model to a decision-grade ~12.6%; a circuit time-constant 1000× off (1 ns vs the correct 1 ps) found and fixed; and the design recorded as FAILING the Shor-2048 budget even at an *idealized, not-achieved* 1:1 ratio (the ~9.6–14.5:1 figure is a *memory-overhead floor*, a different and much smaller quantity than a full-algorithm budget — never equated). All SIMULATED / BOUNDED as marked.

Each of these is a place the engineering corpus made its own claim *weaker* in public. That is the same instrument as this paper's status discipline — the declared weakest links (§0.2), the scoped-out sectors shipped as scoped-out (§2.2, §12), the derived-coupling rows kept as identities and not promoted to predictions (§1.2.1) — and it is why the chamber bridge is offered as evidence rather than decoration: **a program that audits itself this hard, in two independent domains, is showing its work.**

16.4 What §16 does and does not license

It licenses (the honest reading)	It does NOT license (the forbidden over-read)
"The K_6 chamber idea is concrete enough to implement and simulate"	"The QC ratio makes the four-force routing correct" — it does not, and no such inference is drawn; §16 changes no §§5–13 status
"The geometry whose isometries are the forces has a working engineering application at simulated grade"	"The design is a fabricated, measured device" — it is simulated, not built in silicon
"The same self-audit discipline governs both corpora"	"The retired sub-2 ratios are real" — they are withdrawn, and appear nowhere as claims
"An independent domain finds the same coset chamber structure useful"	Any feedback from QC numbers into the interface claims — barred by C-FF13 (no-target-loading)

The standing geometry-change discipline (MR-6, the mirror image of C-FF0's one-parent-geometry rule) is the mirror image of this section: any future geometry move in the corpus must answer *"would it lower the physical:logical qubit ratio at 10,000 logical qubits?"* — a gate that exists **precisely because the geometry is shared**, so a change made to flatter the physics could silently harm the machine, and vice versa. §16 is that gate's positive face: the record that, used forward, the geometry *builds* something. The four-force projections of §§5–9 and the engineering chamber of this section answer to **one** geometry — and that is the strongest reason to take the geometry seriously.

Appendix A — Symbol Glossary

Symbol	Meaning
M_{13}	parent 13D product manifold $M_{3,1} \times K_6 \times S^2 \times S^1$
$M_{3,1}$	4D Lorentzian spacetime factor
$\mathcal{M}_4$	4D spacetime base used in projections (synonymous with $M_{3,1}$)
K_6	six-dimensional internal flag manifold $SU(3)/T^2$
T^2	maximal Cartan torus of $SU(3)$, isomorphic to $U(1)^2$
S^2	two-sphere carrying weak spin-cover structure
S^1	hypercharge circle
S^1_Y	hypercharge circle with boundary parity, identical to S^1 in this paper
G_{MN}	13D metric, primitive bosonic field at the $\times$-layer
$G_{\mu\nu}$	4D spacetime block of the 13D metric
$G_{\mu A}$	mixed metric block whose zero modes descend to 4D gauge connections
G_{AB}	internal metric block carrying compact spectra and moduli
$g_{\mu\nu}$	4D effective metric after compactification
A^a_μ	4D Yang–Mills gauge connection (descended, not primitive)
$SU(3)_c$	surviving color gauge group from K_6
$SU(2)_L$	weak isospin group from S^2 spin cover
$U(1)_Y$	hypercharge group from S^1_Y
$U(1)_{\text{em}}$	unbroken electromagnetic group, $Q = T_3 + Y$
Q	electric-charge operator $T_3 + Y$
T_3	third component of weak isospin
Y	hypercharge
$\mathcal{C}_{\text{admiss}}$	admissibility data at the $\oplus$-layer (selects which configurations count)
F^+_{finite}	finite chamber / discrete projector data at the $\oplus$-layer
$\mathcal{E}_{\text{matter}}$	matter bundle at the $\otimes$-layer
$\mathcal{E}_{\text{gauge}}$	gauge bundle at the $\otimes$-layer
$\mathcal{E}_{\text{Higgs}}$	Higgs / Wilson-line bundle at the $\otimes$-layer
$\mathcal{E}_{\text{proton}}$	proton-stability bundle at the $\otimes$-layer
π_{grav}	projection from the active branch to gravity EFT
π_{strong}	projection from the active branch to QCD EFT
π_{weak}	projection from the active branch to electroweak EFT
π_{QED}	projection from electroweak EFT to QED / electrostatics
$\mathfrak{B}_{\text{active}}$	active branch of the layered architecture, $\times$/$\oplus$/$\otimes$ composition

Appendix B — Force-Routing Tables

Per-force ×/⊕/⊗ routing tables, gathered here for quick reference. Source: Sections 6–9 of this paper.

B.1 Gravity ×/⊕/⊗ Routing

Layer	Gravity content	What must not be smuggled
×	$\mathcal{M}_4$, 13D metric base, $G_{\mu\nu}$, compact geometry	gauge-fixing rules, Hilbert space, ghosts
⊕	gauge choice, diffeomorphism quotient, BRST/BV rules, freeze/downgrade protocol	metric dimensions
⊗	perturbations $h_{\mu\nu}$, stress-energy $T_{\mu\nu}$, ghost/antifield domains, physical Hilbert quotient	base geometry

B.2 Strong ×/⊕/⊗ Routing

Layer	Strong-sector content
×	$K_6 = SU(3)/T^2$, compact spectrum, color geometry
⊕	center/charge rules, regulator, freeze discipline, color-singlet admissibility, θ-scope rule
⊗	quark fields, color modules $V_{SU(3)}$, gluon connection, field strength, BRST/ghost domains

B.3 Weak ×/⊕/⊗ Routing

Layer	Weak-sector content
×	S^2, $S^1_Y/\mathbb{Z}_2$, compact spectra
⊕	chirality/parity rules, $\mathbb{Z}_2$, $\mathbb{Z}_6$, projectors, $Q = T_3 + Y$, freeze protocol
⊗	weak gauge fields, hypercharge field, matter doublets/singlets, Higgs/Wilson-line operator, lepton/quark domains

B.4 QED/Electrostatic ×/⊕/⊗ Routing

Layer	QED/electrostatic content
×	$\mathcal{M}_4$, electroweak compact carriers S^2, S^1_Y
⊕	$Q = T_3 + Y$, $U(1)_{\text{em}}$ selection, gauge-fixing, static-limit rule, conservation rule
⊗	photon field A_μ, charged matter, current J^μ, field strength $F_{\mu\nu}$, stress-energy

Appendix C — No-Double-Counting Ledger

Each row identifies a potential confusion between two parts of the active branch and states the rule that prevents the same object from being counted twice.

Mode / object	Allowed home	Not allowed
$G_{\mu\nu}$ zero mode	gravity	also counted as gauge boson
$G_{\mu A}$ zero modes	gauge-origin bridge	counted as independent primitive Yang-Mills fields
compact scalar modes	modulus / EFT correction	hidden flavor fit knobs
quark zero modes	matter bundle	threshold packet and physical matter simultaneously without state-counting rule
ghosts	gauge-fixing sector	physical matter states
Wilson-line Higgs mode	Higgs sector	generic scalar counterterm
photon	post-EW QED	primitive hypercharge field alone

> **State-counting rule.** A state may enter more than one table only if the tables use different accounting roles. For example, a quark may appear as physical matter in the QCD action and as an anomaly-trace contributor, but it may not be counted as two independent physical states unless a degeneracy factor is explicitly declared.

Appendix D — Provenance Note

This paper was consolidated from earlier project sector notes. Those notes are not required to read or evaluate this paper; the immediate prior draft was *quantum_forces.md*, which this build supersedes (the per-section mapping is in the build header). The only required companion document is Paper I (*Fable_GUT*).

Appendix Q — Interface Certificate Registry

One runnable folder per interface certificate, schema identical to Paper I's R0.R set: `frozen_inputs.yaml`, `src/check.py` (exact arithmetic / structural lint), `run.sh`, `outputs/`, `hashes/`, `falsifiers/`. Each embeds its negative control; the no-doc-only-completion rule applies.

Folder	Checks (one line)	Result	Negative control
`Q01_backbone_coverage/`	All 15 backbone IDs covered in the §5.6 master table (enforcer + certificate + f	PASS	FAIL_CONFIRMED
`Q02_charge_consistency/`	Q = T3 + Y exact on every published charge row	PASS	FAIL_CONFIRMED
`Q03_no_double_counting/`	Appendix C ledger	PASS	FAIL_CONFIRMED
`Q04_composite_qed/`	pi_QED factors through the electroweak EFT	PASS	FAIL_CONFIRMED
`Q05_spec_completeness/`	Every module spec table carries all nine template fields plus the projection row	PASS	FAIL_CONFIRMED
`Q06_sanity_targets/`	The Section 11 table carries the five frozen target formulas verbatim, and each	PASS	FAIL_CONFIRMED
`Q07_scope_vocab/`	Only the four §2.3 status labels are used formally (legacy labels appear solely	PASS	FAIL_CONFIRMED

Coverage is the §5.6 master table's Certificate column: Q01 ↔ C-FF14 (and table completeness), Q02 ↔ C-FF4/C-FF10, Q03 ↔ C-FF1/3/5/6, Q04 ↔ C-FF8, Q05 ↔ C-FF0/2/7, Q06 ↔ C-FF9/11, Q07 ↔ C-FF12/13. Pipeline-level physics regeneration remains Paper I's `reproduce_all.py`; the Q-set (each row PASS above, a runnable computation) certifies this paper's interface claims and published tables at Interface-claim grade.